



Circuits and Systems
Mekelweg 4,
2628 CD Delft
The Netherlands
<http://ens.ewi.tudelft.nl/>

CAS-2021-??

M.Sc. Thesis

Sound Zones with a Cost Function based on Human Hearing

Niels Evert Marinus de Koeijer B.Sc.

Abstract

TODO: Currently, this contains the draft of my thesis.
Everything and anything can be changed as far as I'm
concerned, I appreciate all feedback!

Sound Zones with a Cost Function based on Human Hearing

Subtitle Compulsory?

THESIS

submitted in partial fulfillment of the
requirements for the degree of

MASTER OF SCIENCE

in

ELECTRICAL ENGINEERING

by

Niels Evert Marinus de Koeijer B.Sc.
born in Delft, The Netherlands

This work was performed in:

Circuits and Systems Group
Department of Microelectronics & Computer Engineering
Faculty of Electrical Engineering, Mathematics and Computer Science
Delft University of Technology



Delft University of Technology

Copyright © 2021 Circuits and Systems Group
All rights reserved.

DELFT UNIVERSITY OF TECHNOLOGY
DEPARTMENT OF
MICROELECTRONICS & COMPUTER ENGINEERING

The undersigned hereby certify that they have read and recommend to the Faculty of Electrical Engineering, Mathematics and Computer Science for acceptance a thesis entitled “**Sound Zones with a Cost Function based on Human Hearing**” by **Niels Evert Marinus de Koeijer B.Sc.** in partial fulfillment of the requirements for the degree of **Master of Science**.

Dated: September 15, 2021

Chairman:

dr.ir. R.C. Hendriks

Advisors:

dr. M. Bo Møller

dr. P. Martinez Nuevo

Committee Members:

dr. M. Mastrangeli

dr. J. Martinez Castañeda

Abstract

TODO: Currently, this contains the draft of my thesis. Everything and anything can be changed as far as I'm concerned, I appreciate all feedback!

Acknowledgments

I would like to thank dr. M. Bo Møller, dr. P. Martinez Nuevo, dr.ir. R.C. Hendriks, and dr. J. Martinez Castañeda for all their help and support during the project.

A special thanks to dr. M. Bo Møller for pushing me in the right direction through our numerous white board sessions and to dr. J. Martinez Castañeda for being a major help inside and outside of the project.

Niels Evert Marinus de Koeijer B.Sc.
Delft, The Netherlands
September 15, 2021

Contents

Abstract	v
Acknowledgments	vii
1 Introduction	1
1.1 Preface	1
1.2 Objectives and Organization	3
1.2.1 Answering Research Question 1	3
1.2.2 Answering Research Question 2	3
2 Perceptual Model Review and Implementation	5
2.1 Introduction	5
2.2 Perceptual Background	6
2.2.1 Workings of the Ear	6
2.2.2 Auditory Masking	6
2.2.3 Threshold of Quiet	6
2.3 Review of Perceptual Models	7
2.3.1 Objective Measures	7
2.3.2 Perceptual Models from Audio Coding	11
2.4 Selection of Perceptual Model	14
2.4.1 Criteria for Selecting Perceptual Model	14
2.4.2 Considering Objective Audio Measures	14
2.4.3 Considering Perceptual Models from Audio Coding	15
2.5 Implementation of Selected Perceptual Model	16
2.5.1 Introduction to Par Detectability	16
2.5.2 Computation Details of the Par Detectability	16
2.5.3 Calibration of the the Par Detectability	18
2.5.4 Least-Squares Formulation of the Par Detectability	21
2.6 Conclusion	22
3 Sound Zone Approach Review and Implementation	23
3.1 Introduction	23
3.2 Description of Sound Zone Problem	24
3.3 Sound Zone Data Model	26
3.3.1 Room Topology	26
3.3.2 Defining Target Pressure	26
3.3.3 Realizing Sound Pressure through the Loudspeaker	27
3.3.4 Choice of Target Pressure	28
3.4 Sound Zone Approaches	30
3.4.1 Pressure Matching	31
3.4.2 Acoustic Contrast Control	32
3.5 Selection of Sound Zone Approach	34

3.5.1	Mathematical Properties of Par Detectability	34
3.5.2	Selecting an Appropriate Sound Zone Approach	35
3.6	Conclusion	37
4	Implementation of Perceptual Sound Zone Algorithms	39
4.1	Introduction	39
4.2	Short-Time Reformulation of Pressure Matching	40
4.2.1	Modeling Block Based Knowledge Limitations	40
4.2.2	Block Based Loudspeaker Input Signal Computation	42
4.2.3	Block Based Sound Pressure	42
4.2.4	Derivation of Block-Based Multi-Zone Pressure-Matching	43
4.3	Frequency Domain Reformulation of Pressure Matching	45
4.3.1	Frequency Domain Transformation	45
4.3.2	Quantities the Frequency Domain	46
4.3.3	Proposed Frequency Domain Approach	46
4.4	Unconstrained Perceptual Pressure Matching	48
4.4.1	Summary of Par Detectability	48
4.4.2	Summary of Pressure Matching	48
4.4.3	Detectability Based Pressure Matching	49
4.4.4	Proposed Approach	49
4.5	Constrained Perceptual Pressure Matching	50
4.6	Conclusion	51
5	Evaluation of Perceptual Sound Zone Algorithms	53
5.1	Introduction	53
5.2	Evaluation Methodology	54
5.2.1	Room Setup	54
5.2.2	Content Selection	54
5.2.3	Simulation Outputs	55
5.2.4	Evaluation Criteria	55
5.3	Evaluation of Proposed Algorithms	57
5.4	Conclusion	58
6	Conclusion	59
A	Extra Results	65

Introduction

1.1 Preface

Sound systems are used world wide to fill rooms with enjoyable audio content. Problems arise however when multiple people in the same room want to enjoy different audio content at the same time.

For example, one person may want to enjoy a movie on television, while another may want to listen to their music. If they are in the same room, their desires clash: neither person can fully enjoy their chosen activity without disturbing the other. In short, the interference of multiple source of audio leads to a situation where both individual experiences are diminished.

In recent years, attempts have been made to solve this problem by controlling the spatial reproduction of sound in such a way that different areas in a room have distinct content.

One class of algorithms that attempt to do so are sound zone algorithms [1]. Sound zone algorithms partition the space of the room into multiple so-called sound zones. Each sound zone is assigned different audio content. The sound zone algorithms decides how to use the loudspeakers of the sound system to reproduce said audio content in every zone. Using the principals of constructive and destructive interference, this is done in such that there is minimal interference between zones. That is to say: the audio content of each zone is not audible in the others.

In the previously listed example, one zone would contain the audio of a movie and another zone would contain music. An image depicting the situation is given in Figure 1.1. The sound zone algorithm determines how to best use the sound system to reproduce these two zones. In the ideal case, both people can now enjoy the full potential of their audio content, without bothering one another.

In practice however, the sound zone algorithm will not always do a perfect job. The performance of algorithms depends on the environment and the available sound system. Depending on the situation, the interference between zones can typically only be reduced by so much. As such, audio content of one zone is often still audible in other zones.

Improving sound zone algorithms is thus still an active topic of research. One recent approach is to include a model of the human auditory system, which models how sound is perceived by humans. Typically, sound zone algorithms use sound pressure, which is the physical phenomenon of sound. Sound pressure doesn't always capture what is important for the perception of sound. As such, including a perceptual




Figure 1.1: A room containing a sound system consisting of an array of loudspeakers and two zones. The goal of the sound zone algorithm is to control the sound system in such a way that the red zone contains the audio of a movie, and the blue zone contains the music.

model may allow the algorithm to focus on the parts of the audio content that matter perceptually.

Early results show that the perceptual sound zone approach is promising. Recent work by Donley et al. explored including the absolute threshold of hearing, which models the lowest sound pressure humans can hear, into sound zone algorithms. This pursuit found an increased quality of the reproduced audio in the zones [2]. Other work by Lee et al. showed that including a perceptually-motivated weighting in the sound zone algorithm outperforms traditional algorithms [3, 4].

This work will seek to further explore this perceptual approach by proposing a novel perceptual sound zone algorithm and exploring the possible benefits.

1.2 Objectives and Organization

This section will state the goals of the thesis and organization of the rest of this document.

This thesis seeks to answer two research questions:

- **Research Question 1:** *“How can models of human auditory perception be included in sound zone algorithms?”*
- **Research Question 2:** *“What are the benefits of including models of human auditory perception in sound zone algorithms?”*

What follows is how these research questions will be answered, alongside the structure of the rest of this document.

1.2.1 Answering Research Question 1

The first research question,

*“How can models of human auditory perception be included in
sound zone algorithms?”*

will be answered in chapter 2, chapter 3, and chapter 4. These chapters will document the design of a perceptual sound zone algorithm. The chapters are structured as follows.

- First, in chapter 2
- Next, in chapter 3
- Finally, in chapter 4

1.2.2 Answering Research Question 2

The second research question,

*“What are the benefits of including models of human auditory perception in
sound zone algorithms?”*

will be answered in chapter 5. This will be done by comparing the perceptual sound zone algorithm from Research Question 1 with a reference sound zone algorithm. The chapter is structured as follows.

Perceptual Model Review and Implementation

2

2.1 Introduction

The goal of this chapter is to find a suitable perceptual algorithm for the creation of a perceptual sound zone algorithm. This chapter is structured as follows.

The chapter will begin in section 2.2 with some background information on the human auditory system. This is done to ensure that the reader has the knowledge to understand the perceptual aspects of the thesis.

Next, with the necessary psycho-acoustical background in place, a literature review into perceptual models is given in section 2.3. The purpose of this review is to document candidates for the perceptual model that will be used in the perceptual sound zone algorithm. In addition to this, the reviewed models could also serve as potential candidates for use in the evaluation of the perceptual sound zone algorithm that will be proposed in chapter 4.

To perform the selection of a perceptual model from the candidates discussed in the literature review, criteria reflecting desirable properties for the model for use in sound zones are defined in section 2.4. The criteria are then used to select a perceptual model.

Afterwards, the selected perceptual model is discussed in more detail in section 2.5 by stating implementation details and describing its behavior. The chapter will then wrap up with its conclusions in section 2.6.

2.2 Perceptual Background

This section will provide some additional information on the human auditory system. This serves to give the reader the necessary background information in order to understand the rest of this chapter.

2.2.1 Workings of the Ear

Basilar Membrane, ERBS, ERB

2.2.2 Auditory Masking

Temporal, Spectral

2.2.3 Threshold of Quiet

2.3 Review of Perceptual Models

This section will document a literature review into perceptual models of the human auditory system. As stated in the introduction, this literature review is performed to determine which perceptual models are currently available in the state of the art. This review will then be later used in section 2.4 to select a perceptual model most suitable for integration with a sound zone algorithm.

Especially promising are models that attach some “score” or “rating” to the perceptual quality of input signals. These ratings can then be used in algorithms to obtain an optimal rating through optimization. In addition to this, they can possibly be used to quantify the quality of results from later algorithms.

As such, the focus of the literature review is not on the latest findings in the field of psycho-acoustics, or models that most accurately emulate the behavior of the human ear, but rather on algorithms that quantify quality.

To this end, two categories of perceptual models are considered. First in subsection 2.3.1, “objective measures” are discussed. These are models which attempt to predict the perceptual quality ratings from listening tests. Next, perceptual models from “audio coding” are discussed in subsection 2.3.2. These models are used to quantify how audible the artifacts of audio compression are.

2.3.1 Objective Measures

In order to determine the perceived quality of audio one approach is to use listening tests. These are tests in which subjects are asked to rate a property (or properties) of a set of audio stimuli. One example where listening tests are used is for the evaluation of speech intelligibility of listening aids [5]. Another example is determining which loudspeaker has the best perceived sound quality.

Performing listening tests is however often cumbersome due to the large amount of human labour involved. This motivates the use of objective quality measures, which attempt to predict the outcomes of these listening tests. This is very useful for algorithm developers, as they can get an indication of how well they are doing without having to perform a labour intensive listening test.

Note however that a objective quality measure does not replace a listening test: it can only be used to give an indication. Findings should always be confirmed with listening tests.

The objective measures that will be considered take a reference and degraded audio stimuli as inputs. Most of the discussed models take the following approach. First, they convert the input stimuli to their so-called internal representations, which models how the human auditory system perceives the stimuli. Various features are then derived from the internal representation of the stimuli. The features are then mapped to a prediction of the results of a listening test.

These objective quality measures are promising for integration into sound zone algorithms as they summarize the quality of a signal into a single value, which can

be potentially optimized for. It stands to reason that if an objective quality measure correlates with audio quality, optimizing over such a measure could improve sound zone algorithms.

As such, this section will explore various objective measures. This will be done by considering various classes different objective measures, namely audio quality the quality and intelligibility of speech.

2.3.1.1 Objective Speech Quality Measures

There have been a number of attempts to create objective measures to quantify the perceived quality of speech. In this section three objective speech quality measures will be discussed. Namely the Perceptual Evaluation of Speech Quality (PESQ) [6] measure, Perceptual Objective Listening Quality Assessment (POLQA) [7]. measure, and Virtual Speech Quality Objective Listener (ViSQOL) [8, 9] measure.

- PESQ is a metric which attempts to determine the perceived quality of speech. It was standardized by the International Telecommunication Union (ITU-T) in 2001. PESQ is computed by first applying an auditory transform that maps the reference and degraded speech into a time-frequency representation that models the perceived loudness of the signals. From this internal representation, so-called symmetric and asymmetric disturbances are determined by computing differences between the time-frequency bins of the reference and degraded speech. A non-linear average is then taken to obtain the average disturbance per time bin. These averaged disturbances are then mapped to the outcomes of listening test outcomes through linear combination [6].
- POLQA is speech quality metric which was standardized by the International Telecommunication Union (ITU-T) in 2011. It was intended to be the successor of PESQ, with the improvement of having more accurate predictions on a wider range of distortions. POLQA works with a similar internal representation to PESQ, but computes distortion in a different way as to be capable of handling global temporal compression and expansions [7].
- ViSQOL is a metric developed in 2012 in a collaboration between Trinity College and Google. ViSQOL uses a different internal representation than PESQ and POLQA as it uses neurograms rather than loudness representations. Neurograms contain the neural firing activity of the auditory nerve in time-frequency bins, and NSIM determines how similar the firing patterns of two neurograms are. The neurograms are then compared by means of the Neurogram Similarly Index Measure (NSIM). This similarity is then related to the outcomes of listening tests through a laplacian fit [8], which is then used to make predictions.

In general, PESQ, POLQA and ViSQOL require many steps to compute and were found difficult to optimize for due to conditional branches within the algorithms and many non-differentiable steps such as clipping. Some attempts have been made however to reformulate PESQ in order to make it more tractable for optimization by approximating the disturbances by other functions [10].

2.3.1.2 Objective Speech Intelligibility

Intelligibility of speech is defined as the percentage of words identified correctly given a degraded speech signal. Objective speech intelligibility metrics seek to predict this percentage. In this section, two of these metrics will be discussed. Namely, the Short Time Objective Intelligibility (STOI) [5] measure and the Speech Intelligibility In Bits (SIIB) [11] measure.

- STOI was proposed by Taal et al. in 2011 as a speech intelligibility metric that could make accurate predictions for a speech signals that were degraded time-frequency weighted distortions.

For its internal representation, it finds a time-frequency internal representation through filtering the input stimuli with a filter bank consisting of 1/3 octave bands, and then segmented the filter taps into short time frames. Silent bins that do not contain speech are removed, and clipping is applied to limit the effect of one severely degraded time-frequency bin. The average correlation coefficient between the time-frequency bins of the internal representation of the reference and degraded segments is then computed, and averaged over all bins to determine the intelligibility [5].

- SIIB was introduced by Van Kuyk et al. in 2017 as a speech intelligibility metric that could be motivated through information theory through the mutual information rate. As such, SIIB is given in bits

The idea behind SIIB is that the intelligibility of speech is related to the information shared between intended and degraded speech. SIIB models how transformation of the reference speech signal to the degraded speech signal as a transmission channel. Among other aspects, this transmission channel includes a model of the human auditory system [11]. This communication channel is then used to compute the mutual information rate.

Both STOI and SIIB are difficult to optimize for directly.

In STOI, the removal of silent regions and the clipping operator are non-differentiable operations. Furthermore, the computation of the correlation coefficient is a non-convex function of the degraded speech.

SIIB is in general non-convex and non-differentiable as it uses the Karhunen-Loève transform and a K-nearest neighbor estimator to compute the mutual information rate. However, if the communication channel is approximated as gaussian, the mutual information can be computed in closed form, and SIIB becomes a differentiable measure.

2.3.1.3 Objective Audio Quality Measures

The previous objective quality metrics are both intended for evaluating speech. In this section, a number of objective quality metrics will be discussed that are designed instead for perceived audio quality. Namely, the Perceptual Evaluation of Audio Quality (PEAQ) [12] and ViSQOLAudio [13]. The latter is an adapted version of the ViSQOL speech quality measures.

- PEAQ is an audio quality metric standardized by the International Telecommunication Union (ITU-T). PEAQ estimates a quality grade by first computing an internal representation of the reference and degraded audio signals. This results in a time-frequency representation of the input stimuli from which a number of perceptually relevant features, referred to by PEAQ as Model Output Variables (MOVs), are extracted. An example of these MOVs are the loudness of the noise or the bandwidth of the input stimuli. These MOVs are then mapped to the final audio quality grade through a neural network [12].
- In 2015, it was found that with some adjustments ViSQOL could be used to determine audio quality, which resulted in a new metric ViSQOLAudio. Among the adjustments were the removal of the voice activity detector included in ViSQOL and the use of a larger bandwidth to cover the entire spectrum of hearing from 50 Hz to 20000 Hz, rather than just the bandwidth of speech.

PEAQ, and ViSQOLAudio are both difficult to optimize for. A number of the MOVs computed in PEAQ, such as the partial noise loudness, are non-differentiable. As ViSQOLAudio is similar to ViSQOL with some small adjustments, it is similarly difficult to optimize for.

2.3.1.4 Distraction Model

One especially promising objective measure is the distraction proposed by Francombe et al. in 2015 [14]. This measure was designed with the application of sound zones in mind.

The distraction was determined to be the keyword that best describes the perceptual experience of interfering audio programs. This was determined through an elicitation study performed also performed by Francombe et al. in 2014 [15]. This prompted the creation of the model.

To create the model, a listening test was performed where the participants were subjected to audio-on-audio interference. The subjects were played a target audio stimuli they were instructed to focus listening to. At the same time, an interferer audio stimuli was played to distract the participant from the target. The participants were given a scale between 0 and 100 on which they were asked to rate how distracting the interference was when listening to the target program, where a 100 was maximally distracted.

The target-interferer stimuli pairs and corresponding ratings resulted in a dataset. This dataset was then used to fit a model which predicted the distraction given novel a target-interferer stimuli pair. The model consisted of taking a linear combination of 5 features that were computed from the stimuli.

Computing said features could however not be performed in real time. The reason for this was that as the original distraction model is too computationally complex [16]. To this end, in 2017, Rämö et al. proposed a version of the distraction model that could be run in real-time. This was done by approximating the features of the original distraction model by computationally less complex alternatives. The

resulting real-time distraction model was found to be less precise, but could be run in 0.04% of the time of the original distraction model.

On face value, the real-time distraction model seems promising to optimize over. However, while easy to compute, the model is non-differentiable as the model uses piecewise functions and non-convex due to taking the logarithm of the square of the input signals. In addition to this, the model also performs operations that are difficult to express mathematically, such as counting the number of short-time blocks that exceed a certain threshold.

2.3.2 Perceptual Models from Audio Coding

The second class of perceptual models that will be considered are the perceptual models used in audio coding. Audio coding algorithms attempt to find an low-bitrate representation of an audio input signal, which is a form of lossy compression. As such, reducing the bitrate introduces errors. These errors can be a detriment to the listening experience.

As such, most audio coding algorithms use a perceptual model to quantify how disturbing the introduced distortions are. The perceptual model is used to introduce encoding errors in such a way that the audio output signal is perceptually indistinguishable from the audio input signal [17]. The perceptual model typically takes form of a distortion function which determines how audible the difference between a reference input audio signal and a distorted output audio signal is. This function can be used to for example encode an input audio signal such that it has minimal distortion for a specified bitrate.

The perceptual models used in audio coding are promising for integration into a sound zone algorithm, as they are often mathematically tractable. As stated, these perceptual models typically take the form of some sort of distortion function that quantifies how perceptually disturbing the introduced artifacts are. One approach could be to define sound zone algorithms that minimize a distortion function for example.

As such this section will explore three perceptual models from audio coding.

2.3.2.1 ISO MPEG Models

The ISO/IEC 11172-3 standard specifies a coded representation for audio files [18], and a corresponding decoder. An encoder said representation is not part of the standard. This is done deliberately, to allow for future improvements to the encoder, without having to change the standard [19].

The standard does however provide a number of examples of possible encoders, with increasing complexity. Alongside these example encoders, two psycho-acoustical models are included for use during the encoding process.

The psycho-acoustical models work by subdividing the input audio signal into frequency bands which correspond to the frequency bands in the human auditory system. The model then determines how much quantization noise can be added

separately per band without the noise becoming audible. As such, the model assumes that the distortion signal is noise-like [20], which is usually the case for quantization noise for audio coders.

The output of the psycho-acoustical model is thus the amount of noise that can be added per band. In the case of audio coding, this can then be used to control quantization noise. Note that this perceptual model does not come in the form of the earlier described distortion function. This technique has however been used for various signal processing purposes, such as audio watermarking [17]. As such, examples exist from which optimization schemes could be inspired.

2.3.2.2 Par Detectability

In 2005, van der Par et al. proposed a novel perceptual model designed for use in audio coding [20]. The model defines a distortion measure which determines the “detectability” of a distortion signal in presence of masking signal. That is to say, the function quantifies the degree to which a human is to detect a distortion signal. For audio coding purposes, this distortion signal is error introduced due to the audio compression.

Similarly to the ISO MPEG perceptual models, the Par detectability typically operates on short-time segments, typically in the order of 20 to 200 milliseconds. The proposed method however differentiates itself from the previously discussed ISO MPEG models in three ways.

Firstly, the paper uses newer findings from psycho-acoustic literature, namely spectral integration. In spectral integration, the masking effects from neighboring bands are taken into account when computing the masking effects. The psycho-acoustical models defined in the ISO MPEG standard does not do this as it effectively works independently per band [17].

Secondly, it assumes that the distortion signal is sinusoidal, rather than noise-like. As such, it is more effective in hiding sinusoidal distortion.

Thirdly and finally, the perceptual model is described as a distortion function which quantifies how detectable a disturbance stimuli is. The proposed distortion measure can be expressed as an L2-norm.

This mathematical tractability makes for easy integration into existing least-square problems. As such, the Par model has been used in many signal processing applications, examples ranging from speech enhancement to removing perceptually irrelevant sinusoidal components [21, 22].

2.3.2.3 Taal Detectability

A paper from 2012 by Taal et al. proposed a novel perceptual model [17] which introduces an alternative definition to the detectability defined in the Par model.

In contrast to the Par detectability, the Taal detectability measure takes temporal characteristics of a signal into account. The inclusion of temporal information allows for the suppression of “pre-echoes”, which is an artifact that the Par model suffers

from. The “pre-echoes” artifacts arises from the assumption that the masking effects of the masking signal are stationary across time. As a result, audio coding algorithms may assume that audio content is masked while it is not, which could results in quantization noise not being masked.

In contrast to other temporal perceptual models, the Taal Detectability has a relatively low computational complexity. In addition to this, it can also be expressed as an L2-norm. The computational demand was however shown to be higher than the Par Detectability [17], especially for longer time segments.

2.4 Selection of Perceptual Model

In section 2.3 a literature review of various models of the human auditory system was given. This section will detail the selection of one of these perceptual models for use in a perceptual sound zone algorithm.

The structure of this section is as follows. First, subsection 2.4.1 will discuss the criteria that will be used to select the most suitable perceptual model for use in a perceptual sound zone algorithm. Next, these criteria will be used to evaluate how promising the candidate models are. This will be done by first considering the objective audio measures in subsection 2.4.2, followed by the perceptual models from audio coding in subsection 2.4.3.

2.4.1 Criteria for Selecting Perceptual Model

This section will define desirable criteria for the perceptual model for integration with a sound zone algorithm.

The main criteria is mathematically tractability. That is to say, it is desirable for the perceptual model to be easy to include mathematically in optimization problems.

As will be shown in chapter 3, many sound zone algorithms are posed as optimization problems. Therefore, it is undesirable if a perceptual model is too complicated to be integrated mathematically into the existing optimization problems. For example, if the computation of a perceptual model involves heavy conditional branching (i.e. if-statements), their integration into an existing sound zone optimization problem may be challenging.

Furthermore, most sound zone algorithms can be posed as a convex optimization problem. Convexity is a desirable property for an optimization problem, as it guarantees that the optimization problem has a single, global optimal value, rather than many locally optimal values [23]. In addition to this, there are many efficient solvers available for convex optimization problems. Therefore, it is also desirable that the perceptual models preserve this convexity.

2.4.2 Considering Objective Audio Measures

All objective audio measures were found to be mathematically untractable. Every model showed a degree of non-differentiability and non-convexity in their computation. For some models, such as the distraction model, the complete computation could not be expressed easily with mathematical expressions. As such, they are difficult to integrate into convex optimization problems and will not be used in the perceptual sound zone algorithm. However, as the objective audio measures predict the outcomes of listening tests, they will be used for evaluation.

2.4.3 Considering Perceptual Models from Audio Coding

In contrast to the objective audio measures, all models from audio coding are mathematically tractable.

The ISO MPEG models were found to be less promising than the Par and Taal detectability, as it does not immediately define a cost function which can be optimized over: the models only determine how much noise can be added per band.

As such, the decision is between the Par and Taal detectability. The Taal detectability takes into account temporal properties of the input signal in its perceptual model. This is beneficial, as it will lead to a more accurate description of the masking properties of the input signals. However, it has been shown to be at the cost of computational complexity.

Surface level inspection has shown Taal and Par detectability seem sufficiently similar that it is likely possible to use the Taal detectability in place of the Par detectability in the algorithms proposed in chapter 2. The Taal detectability is evaluated in the time domain rather than the frequency domain, which is the case for the Par model. Exploring the possibilities of the Taal detectability will however be left to future work and not further explored in this work.

As such, the Par detectability will be selected for use in the sound zone algorithm.

2.5 Implementation of Selected Perceptual Model

In section 2.4, it was determined that the Par detectability was the most suited model for the perceptual sound zone algorithm of all perceptual models considered in the literature review given in section 2.3. In this section, the Par detectability is considered in greater detail in order to give the reader a greater understanding of the model.

This section is organized as follows. First, subsection 2.5.1 gives a high-level description of the detectability, providing an intuitive understanding and introducing some of the notation that will be used. Next, the steps to computing the detectability are described in subsection 2.5.2. Afterwards, subsection 2.5.3 discusses the calibration of the Par detectability. Finally, subsection 2.5.4 rewrites the detectability to a more mathematically tractable expression.

2.5.1 Introduction to Par Detectability

In this section, a high-level description of the Par detectability will be given. This is done to give the reader a basic understanding of the model before going into greater detail. The Par detectability maps two input sequences to a positive real value, i.e. $D : (\mathbb{R}^{N_x}, \mathbb{R}^{N_x}) \mapsto \mathbb{R}^+$. The two input sequences will be denoted as the masking signal $x[n] \in \mathbb{R}^{N_x}$ the disturbance signal $\varepsilon[n] \in \mathbb{R}^{N_x}$. The detectability of these two sequences will be denoted as $D(x[n], \varepsilon[n])$.

It is assumed that a human is listening to both the masking signal $x[n]$ and the disturbance signal $\varepsilon[n]$ at the same time. The detectability $D(x[n], \varepsilon[n])$ can then be understood as how easily a human listener can detect the disturbance signal $\varepsilon[n]$ in presence of the masking signal $x[n]$. The signal $x[n]$ is referred to as the masking signal because it is assumed to mask the disturbance signal $\varepsilon[n]$ to a degree.

For this interpretation to be accurate, the signals $x[n]$ and $\varepsilon[n]$ should be short-time signals. The paper uses a signal length N_x corresponding to a length in time between 20 to 200 milliseconds. This is important, as the model assumes that the psycho-acoustical properties of $x[n]$ and $\varepsilon[n]$ are stationary.

The metric is normalized in such a way that the detectability $D(x[n], \varepsilon[n])$ is equal to 1 when the disturbance signal $\varepsilon[n]$ is “just noticeable” in presence of masking signal $x[n]$. That is to say: if the detectability is 1, the disturbance is on the verge of being noticeable and not noticeable.

The detectability $D(x[n], \varepsilon[n])$ can also attain a value larger than 1. The larger values of the detectability correspond with an increased perceived presence of the disturbance signal $\varepsilon[n]$.

2.5.2 Computation Details of the Par Detectability

This section will explore calculating the Par detectability. The first thing to note about the Par detectability is that it is computed using the frequency domain representations of its inputs. To this end, let $X[k]$ and $\mathcal{E}[k]$ denote the frequency

domain representations of the masking signal $x[n]$ and the disturbance signal $\varepsilon[n]$ respectively.

After determining the frequency domain representations, the Par detectability computes an internal representation of the input signals $X[k]$ and $\mathcal{E}[k]$. This internal representation models how the input stimuli appear to the human auditory system, and is modeled by applying a filter to the input stimuli.

Two subsequent filters are applied. The first filter models how parts of the ear filter the incoming sound with an outer- and middle-ear filter $H_{\text{om}}[k]$. Next, a 4th order Gammatone filter bank is applied, modeling the frequency-place transform that occurs in basilar membrane inside of the ear [20].

The Gammatone filter bank consists of N_g filters. The frequency domain representation of each individual filter is denoted by $\Gamma_i[k]$, for $1 \leq i \leq N_g$. The filters in the filter bank $\Gamma_i[k]$ have a bandwidth given by the equivalent rectangular bandwidth (ERB) and center frequencies given by the corresponding equivalent rectangular bandwidth number scale (ERBS)¹.

After filtering, the power per Gammatone filter tap is computed. Let M_i and S_i denote the output power of the i^{th} filter tap for the masking signal $X[k]$ and the disturbance signal $\mathcal{E}[k]$ respectively. The relationship between the input quantities and the output power of the filter taps can be given as follows:

$$M_i = \frac{1}{N_x} \sum_{k=0}^{N_x-1} |H_{\text{om}}[k]|^2 |\Gamma_i[k]|^2 |X[k]|^2 \quad (2.1)$$

$$S_i = \frac{1}{N_x} \sum_{k=0}^{N_x-1} |H_{\text{om}}[k]|^2 |\Gamma_i[k]|^2 |\mathcal{E}[k]|^2 \quad (2.2)$$

The output powers can then be used to define the within-channel detectability D_i per filter tap i . This can be thought of the detectability per filter tap. The within-channel detectability is defined as follows:

$$D_i = \frac{N_x S_i}{N_x M_i + C_a} \quad (2.3)$$

Here, C_a is a calibration constant that ensures that the absolute threshold of hearing is predicted correctly. This can be understood by considering the case where no masking signal $x[n]$ is present, in which case $M_i = 0$ for all i . If not for the calibration constant C_a , the detectability of any non-zero disturbance $\varepsilon[n]$ would be infinite.

The total detectability $D(x[n], \varepsilon[n])$ can then be computed as the scaled sum of all

¹Expressions for the gammatone filters $\Gamma_i[k]$ are provided by the original paper [20].

within channel detectabilities. It is defined as follows:

$$D(x[n], \varepsilon[n]) = C_s L_{\text{eff}} \sum_{i=0}^{N_g} D_i \quad (2.4)$$

$$= C_s L_{\text{eff}} \sum_{i=0}^{N_g} \frac{\sum_{k=0}^{N_x-1} |H_{\text{om}}[k]|^2 |\Gamma_i[k]|^2 |\mathcal{E}[k]|^2}{\sum_{k=0}^{N_x-1} |H_{\text{om}}[k]|^2 |\Gamma_i[k]|^2 |X[k]|^2 + C_a} \quad (2.5)$$

Here, C_s is a calibration constant chosen such that a just noticeable disturbance signal results in a detectability of $D(x[n], \varepsilon[n]) = 1$. The constant L_{eff} is the integration time of the human auditory system. It is chosen equal to the segment length of $x[n]$ and $\varepsilon[n]$ in milliseconds.

In order to further understand detectability, consider the behavior of the expression of the detectability $D(x[n], \varepsilon[n])$ above. Imagine that the spectrum of the masking signal is much larger than the disturbance signal, i.e. $X[k] \gg \mathcal{E}[k]$ for all frequency bins k . In this case, the detectability of $\varepsilon[n]$ will be small due to the masking of the masking signal $x[n]$ or inaudible due to the threshold of hearing determined by the calibration constant C_a .

Conversely consider the case that the spectrum of the masking signal is much smaller than the disturbance signal, i.e. $X[k] \ll \mathcal{E}[k]$ for all frequency bins k . In this case, the resulting detectability is determined greatly by the calibration coefficient C_a . If the total energy of the filtered disturbance signal is much larger than the calibration constant $S_i \gg C_a$ for all i , the detectability becomes large. This models the case that the disturbance signal is large relative to the threshold of hearing. Alternatively, if $S_i \ll C_a$ for all i , the disturbance signal is inaudible due to the threshold of hearing and the detectability will be low accordingly.

This concludes the analysis of the Par model. What follows is the discussion of the calibration of the Par model in subsection 2.5.3, where calibration constants C_a and C_s are determined. Afterwards, the model above will be given as a least-squares formulation in order to ease the integration into optimization in subsection 2.5.4.

2.5.3 Calibration of the the Par Detectability

A calibration is necessary for the Par detectability to provide the correct output. In the previous section, it was shown that the constants C_a and C_s are used to this end. A correct calibration of the Par detectability must satisfy the following:

1. The just noticeable disturbance signal must result in a detectability of 1.
2. The threshold of hearing takes effect appropriately.

Both the concepts of just noticeable distortion and the threshold of hearing require knowledge of the sound pressure level of the stimuli. Thus, before determining the calibration coefficients, it is important to first discuss the relationship between the input signals $x[n]$ and $\varepsilon[n]$ and reproduced sound pressure level. This is the topic of subsubsection 2.5.3.1. Afterwards, subsubsection 2.5.3.2 discusses the determination of the coefficients C_a and C_s .

2.5.3.1 Relating Digital Representation and Sound Pressure Level

One difficulty of taking the threshold of hearing into account is that it is typically given in terms of sound pressure level (SPL), measured in dB. The one-sided spectrum of the threshold of hearing in dB SPL can be approximated by the following function [24]:

$$T_q(f) = 3.64 \left(\frac{f}{1000} \right)^{-0.8} + 0.001 \left(\frac{f}{1000} \right)^4 - 6.5 \exp \left[-0.6 \left(\frac{f}{1000} - 3.3 \right)^2 \right] \quad (2.6)$$

The signals $x[n]$ and $\varepsilon[n]$ are however given digital representation of audio.

For example, they might be given in a pulse code modulated (PCM) format within which they attain integer values between -32786 and 32787.

As such, to meaningfully integrate the threshold of quiet, the digital representation and the sound pressure levels must be related. This relationship can be modeled as follows:

$$X_{\text{dB}}(f) = 10 \log_{10}(|X(f)|^2) + O_{\text{dB}} \quad (2.7)$$

Here, $X_{\text{dB}}(f)$ is the dB SPL representation of a given spectrum $X(f)$. Furthermore, O_{dB} is an offset to ensure the digital representation corresponds to the correct sound pressure level. In order to use this relationship to determine the appropriate digital equivalent of the threshold in quiet, a definition of the offset O_{dB} must be determined.

One way of determining the offset O_{dB} is by relating the sound pressure level and the digital representation of a full-scale sinusoid. A full-scale sinusoid is a sinusoid that has an amplitude of the maximum value that can be attained in the digital representation.

In our previous example, one way of doing so would be to state that a full-scale sinusoid with amplitude 32787 corresponds to e.g. a sound pressure level of 100 dB SPL. The interpretation of this is that playing a full-scale sinusoid will result in a sound pressure of 100 dB SPL when played from the sound system.

To do so, let the digital representation of the full-scale sinusoid be modeled by a sinusoid with amplitude A and frequency f_0 . Consider the one-sided fourier representation of the digital representation of this full-scale sinusoid:

$$\mathcal{F} \{ A \cos(2\pi f_0 t) \} = A \delta(f - f_0) \quad (2.8)$$

It is assumed that playing the digital representation of this sinusoid results in a sound pressure level of A_{dB} db SPL. Substituting these definitions into Equation 2.7 results in the following definition for O_{dB} :

$$O_{\text{dB}} = 10 \log_{10}(|A|^2) - A_{\text{dB}} \quad (2.9)$$

The offset fully defines the relationship between digital representation and sound pressure level, and allows for the conversion of the threshold of hearing to digital representation.

2.5.3.2 Determining Calibration Constants

There are various ways of calibrating this model, but this section will discuss the method of calibrating that is given in the original paper [20]. The given approach is to find the two unknowns C_a and C_s by solving a system of two equations that model the previously stated calibration requirements.

The first requirement is that a just noticeable disturbance signal must result in a detectability of 1. From perceptual literature it is known that a sinusoidal disturbance signal at a given frequency f_0 is just noticeable in presence of an in-phase sinusoidal masking signal that is 18 dB SPL louder [20]. To model this, consider the following masking and disturbance signals.

$$x_{\text{JND}}[n] = A_{70} \cos(2\pi f_0 n / f_s) \quad (2.10)$$

$$\varepsilon_{\text{JND}}[n] = A_{52} \cos(2\pi f_0 n / f_s) \quad (2.11)$$

Here, $x_{\text{JND}}[n]$ is a sinusoid with an amplitude A_{70} , which corresponds to 70 dB SPL. Furthermore, $\varepsilon_{\text{JND}}[n]$ is a sinusoid with an amplitude A_{52} , which is 18 dB SPL less. Note that the amplitudes are both given in digital representation, not sound pressure level representation. The digital representation amplitudes are found through Equation 2.7.

Thus, $\varepsilon_{\text{JND}}[n]$ must be just noticeable in presence of $x_{\text{JND}}[n]$. This can be expressed as follows:

$$D(x_{\text{JND}}[n], \varepsilon_{\text{JND}}[n]) = 1 \quad (2.12)$$

This expression forms the first equation in the system of equations that can be solved to calibrate the Par detectability.

The second requirement is that the threshold of hearing must be included correctly. The threshold of hearing defines the sound pressure levels that are the verge between audible and inaudible sound as a function of frequency. To this end, consider the following masking and disturbance signals:

$$x_{\text{THR}}[n] = 0 \quad (2.13)$$

$$\varepsilon_{\text{THR}}[n] = A_{\text{tq}} \cos(2\pi f_0 n / f_s) \quad (2.14)$$

Here the masking signal $x_{\text{THR}}[n]$ is zero. The disturbance signal is a sinusoid of frequency f_0 with amplitude A_{tq} , which is chosen such that it attains the threshold of quiet at f_0 , i.e. $T_q(f_0)$.

As the threshold of quiet is the verge between audible and inaudible sound, it is assumed that a disturbance signal in presence of no masking signal that has an amplitude equal to the threshold of quiet is just noticeable. Recall that for just noticeable stimuli the detectability must be equal to 1. This allows us to specify the second equation in the system of equations:

$$D(0, \varepsilon_{\text{THR}}[n]) = 1 \quad (2.15)$$

The system of equations defined by Equation 2.12 and Equation 2.15 can be solved through the bisection method. To see how this is done, the reader is referred to the original paper [20].

2.5.4 Least-Squares Formulation of the Par Detectability

This section will rewrite the previously introduced detectability into a least-squares representation. This representation is more mathematically tractable than Equation 2.5 and thus will allow for easier integration into existing sound zone algorithms.

To obtain this expression, the sum of squares will be expressed as a L2-norm. Consider the following rewrite of the detectability given in Equation 2.5:

$$\begin{aligned}
D(x[n], \varepsilon[n]) &= C_s L_{\text{eff}} \sum_{i=0}^{N_g} \frac{\sum_{k=0}^{N_x-1} |H_{\text{om}}[k]|^2 |\Gamma_i[k]|^2 |\mathcal{E}[k]|^2}{\sum_{k=0}^{N_x-1} |H_{\text{om}}[k]|^2 |\Gamma_i[k]|^2 |X[k]|^2 + C_a} \\
&= \sum_{i=0}^{N_g} \left(\frac{C_s L_{\text{eff}}}{\|H_{\text{om}}[k] \Gamma_i[k] X[k]\|_2^2 + C_a} \right) \sum_{k=0}^{N_x-1} |H_{\text{om}}[k]|^2 |\Gamma_i[k]|^2 |\mathcal{E}[k]|^2 \\
&= \sum_{k=0}^{N_x-1} \left(\sum_{i=0}^{N_g} \frac{C_s L_{\text{eff}}}{\|H_{\text{om}}[k] \Gamma_i[k] X[k]\|_2^2 + C_a} \right) |H_{\text{om}}[k]|^2 |\Gamma_i[k]|^2 |\mathcal{E}[k]|^2 \\
&= \sum_{k=0}^{N_x-1} |W_x[k]|^2 |\mathcal{E}[k]|^2 \\
&= \|W_x[k] \mathcal{E}[k]\|_2^2
\end{aligned}$$

The rewrite above introduced perceptual weighting $W_x[k] \in \mathbb{R}^{N_x}$ describing the masking effects of the masking signal $x[n]$. The perceptual weighting $W_x[k]$ is defined as follows:

$$W_x[k] = \left(\sqrt{\sum_{i=0}^{N_g} \frac{C_s L_{\text{eff}}}{\|H_{\text{om}}[k] \Gamma_i[k] X[k]\|_2^2 + C_a}} \right) |H_{\text{om}}[k]| |\Gamma_i[k]| \quad (2.16)$$

Note from this formulation that the perceptual weighting is only a function of the masking signal $x[n]$. Note also that the resulting detectability $D(x[n], \varepsilon[n])$ is a convex function of the disturbance signal $\varepsilon[n]$.

2.6 Conclusion

The goal of this chapter was to determine a suitable perceptual model for integration into a sound zone algorithm. This was done by first reviewing various objective quality measures and perceptual models from audio coding.

To select a suitable model, the mathematical tractability and the additional computational overhead of each model under review was considered. From this, it was found that the Par detectability was found to be the most promising perceptual model included in the review. Finally, the chapter concluded with an analysis and implementation of the Par detectability.

Sound Zone Approach Review and Implementation

3

3.1 Introduction

The goal of this chapter is find a suitable sound zone approach for integration with the perceptual model selected in chapter 2.

Before selecting an approach, a review of sound zone literature is given. As such, the chapter will start with a brief description of sound zones in section 3.2 to give the reader the necessary background information on the sound zone problem. Afterwards, section 3.3 will present the mathematical framework used to describe the sound zone problem.

This mathematical framework is then used to describe the two main sound zone approaches, “Pressure Matching” and “Acoustic Contrast Control”, in section 3.4. One of these approaches will be used in the perceptual sound zone algorithm.

In order to determine which one is most suited, section 3.5 will reflect on the mathematical properties of the Par detectability selected in chapter 2. The selected approach is then adjusted in ?? and ?? as to prepare it for integration with the perceptual model in chapter 4.

The chapter ends with a summary and concluding remarks in section 3.6.

3.2 Description of Sound Zone Problem

An initial description of the sound zone problem was given in the introduction of the thesis. This section seeks to build on this description in order to provide the understanding necessary for the rest of this work.

As mentioned in the introduction, the problem that sound zones seeks to solve is the reproduction of multiple types of audio content in the same room with minimal interference. This way, multiple people can enjoy different audio content without disturbing one another.

Controlling the spatial distribution of sound is done by controlling the audio that is produced by an array of loudspeakers. The space inside the enclosure is divided up into multiple zones. Each zone is assigned target sound pressure that we would like to have reproduced inside of it. This target sound pressure could be various audio content, for example music or the sound of a movie.

To understand this principal, consider the example given by Figure 3.1.




Figure 3.1: A room is divided into two zones: a red zone and a blue zone. Each zone is assigned different content: content A and content B. The loudspeakers array that is present in the room is to be controlled by the sound zone algorithm in such a way that the desired content is reproduced in each zone. As mentioned, this is to be done in a way that results in minimal interference, e.g. it is undesirable to be able to hear content B when inside the red zone.

There are various approaches to solving the sound zone problem. An important concept that is used frequently is that bright zones and dark zones.

Sound zone problems are typically decomposed into a separate subproblem for every zone. Each one of these subproblems considers only two zones: one bright zone and one dark zone. The goal of each subproblems is to reproduce a specified target sound pressure in the bright zone while restricting the sound pressure in the dark zones. In doing so, the target sound pressure is essentially reproduced locally in the bright zone, while the sound pressure leaking into the dark zone is kept low.

The combination of all subproblems provides a solution to the sound zone problem. To ease the understanding of this concept, consider an example of this decomposi-




Figure 3.2: Decomposition of the example given in Figure 3.1 into two bright-dark zone pairs. For the first problem, the goal is to reproduce content A in bright zone A while minimizing the amount of sound pressure in dark zone A. Similarly for the second problem: reproduce content B in bright zone B while minimizing the amount of sound pressure in dark zone B. Combining the two solutions results in a solution with content reproduced in both zones with minimal interference between zones.

tion is given in Figure 3.2.

This concludes the introduction to the sound zone problem. To further understanding, in the next section, the sound zone problem will be formulated mathematically by introducing the data model.

3.3 Sound Zone Data Model

In the previous section, the sound zone problem was introduced heuristically. In this section a mathematical framework for a room containing sound zones will be introduced. This framework will be used later in the derivation of the sound zone algorithms in section 3.4.

The contents of this section are as follows. First, subsection 3.3.1 develops a spatial description of a room containing two zones and a loudspeaker array. Then, subsection 3.3.2 defines the objective of the sound zone algorithm formally as realizing target sound pressure at discrete points in the room.

The relation between the sound pressure in the room and loudspeaker input signals will then be given in subsection 3.3.3, completing the mathematical framework. This is then used in subsection 3.3.4 to select a suitable target sound pressure which will be used in the remainder of this thesis.

3.3.1 Room Topology

In this section, a description of the room in which sound zones are to be reproduced will be given. In general, the room can contain any number of zones, but this thesis will focus on the two zone case.

The room \mathcal{R} can be modeled as a closed subset of three dimensional space, $\mathcal{R} \subset \mathbb{R}^3$. The two non-overlapping zones \mathcal{A} and \mathcal{B} are contained within the room \mathcal{R} , i.e. $\mathcal{A} \subset \mathcal{R}$ and $\mathcal{B} \subset \mathcal{R}$ where $\mathcal{A} \cap \mathcal{B} = \emptyset$. In addition to the zones, the room \mathcal{R} also contains N_L loudspeakers, whose locations are modeled as discrete points. The room, loudspeakers and zones are visualized in Figure 3.3.

The goal of the sound zone algorithm is to use the loudspeakers to realise a specified target sound pressure in the space described by zones \mathcal{A} and \mathcal{B} . This is to be done in such a way that there is minimal interference between zones, meaning that target sound pressure intended for one zone should not be audible in the other zones.

The loudspeakers can be controlled by specifying their input signals. As such, the goal of the sound zone algorithm is finding loudspeaker input signals in such a way that specified target sound pressure is attained.

The rest of this section will focus on formalizing this notion mathematically.

3.3.2 Defining Target Pressure

As mentioned, the goal of the sound zone algorithm is to realize a specified target sound pressure in the different zones \mathcal{A} and \mathcal{B} in the room \mathcal{R} .

Currently, the zones are given as continuous regions in space. However, most sound zone approaches will instead discretize the zones by sampling the continuous zones \mathcal{A} and \mathcal{B} into so-called control points. The sound pressure is then controlled only in these control points.




Figure 3.3: The room $\mathcal{R} \subset \mathbb{R}^3$ containing the zones $\mathcal{A} \subset \mathcal{R}$ and $\mathcal{B} \subset \mathcal{R}$ depicted in green and blue respectively. The room contains $N_L = 8$ loudspeakers, which are denoted by the red dots in the corners of the room.

Thus, we discretize zones \mathcal{A} and \mathcal{B} into a total of N_a and N_b control points respectively. Let A and B denote the sets of the resulting control points contained within zones \mathcal{A} and \mathcal{B} respectively.

Now let $t^m[n]$ denote the target sound pressure at control point m in either A or B , i.e. $m \in A \cup B$. Our goal is thus to realize $t^m[n]$ in all control points $m \in A \cup B$ using the loudspeakers present in the room. The relationship between the loudspeaker input signals and the sound pressure is the topic of the next section.

3.3.3 Realizing Sound Pressure through the Loudspeaker

The sound pressure produced by the loudspeakers can be controlled by specifying their input signals. Mathematically speaking, let $x^{(l)}[n] \in \mathbb{R}^{N_x}$ denote the loudspeaker input signal of length N_x for the l^{th} loudspeaker. For now, it is assumed that the loudspeaker input signals are of finite length. In a later part of the thesis, a short-time formulation will be given which supports infinite length sequences.

As such, the goal of the sound zone algorithm is to find loudspeaker inputs $x^{(l)}[n]$ such that the target sound pressure $t^m[n]$ is realized for all $m \in A \cup B$.

In order to do so, a relationship must be established between the loudspeaker inputs $x^{(l)}[n]$ and the resulting sound pressure at control points $m \in A \cup B$. This relationship can be modeled by room impulse responses (RIRs) $h^{(l,m)}[n] \in \mathbb{R}^{N_h}$.

The RIRs $h^{(l,m)}[n]$ determine the sound pressure at control point m due to playing loudspeaker signal $x^{(l)}[n]$ from loudspeaker l . Mathematically, let $p^{(l,m)}[n] \in$




Figure 3.4: The previously introduced room \mathcal{R} with zones \mathcal{A} and \mathcal{B} discretized.

$\mathbb{R}^{N_x+N_h-1}$ represent said sound pressure. It can be defined as follows:

$$p^{(l,m)}[n] = (h^{(l,m)} * x^{(l)})[n] \quad (3.1)$$

Here, the $(*)$ operator is used to denote linear convolution. The realized sound pressure $p^{(l,m)}[n]$ only considers the contribution of loudspeaker l at reproduction point m . Let $p^{(l)}[n] \in \mathbb{R}^{N_x+N_h-1}$ denote the total sound pressure due to all N_L loudspeakers, which can be expressed as the sum over all contributions $p^{(l,m)}[n]$ as follows:

$$p^{(m)}[n] = \sum_{l=0}^{N_L} p^{(l,m)}[n] \quad (3.2)$$

$$= \sum_{l=0}^{N_L} (h^{(l,m)} * x^{(l)})[n] \quad (3.3)$$

With the data model completed, the goal of the sound zone algorithm can be restated formally. Namely, the goal is to find $x^{(l)}[n]$ such that the realized sound pressure $p^{(m)}[n]$ attains the target sound pressure $t^{(m)}[n]$ for all control points $m \in A \cup B$.

3.3.4 Choice of Target Pressure

The target sound pressure $t^{(m)}[n]$ describes the desired content for a specific control point m . So far, the choice of target sound pressure $t^{(m)}[n]$ has been kept general. In this section, a choice for the target pressure will be made and motivated.

Assume that the users of the sound zone system have selected desired playback audio signals $s_{\mathcal{A}}[n] \in \mathbb{R}^{N_x}$ and $s_{\mathcal{B}}[n] \in \mathbb{R}^{N_x}$ that they wish to hear in zone \mathcal{A} and

\mathcal{B} respectively. In order to accommodate the wishes of the user, the target sound pressure is chosen as follows:

$$t^{(m)}[n] = \sum_{l=0}^{N_L} (h^{(l,m)} * s_{\mathcal{A}})[n] \quad \forall m \in A \quad (3.4)$$

$$t^{(m)}[n] = \sum_{l=0}^{N_L} (h^{(l,m)} * s_{\mathcal{B}})[n] \quad \forall m \in B \quad (3.5)$$

This choice for the target pressure can be understood as the sound pressure that arises in a certain zone when playing only the desired playback audio for that zone from the loudspeaker array. For example, when in zone $m \in A$, the target sound pressure is set equal to the sound pressure corresponding to the sound pressure that arises when playing only $s_{\mathcal{A}}[n]$ from the loudspeaker array.

The motivation for choosing this target is that it is physically attainable with the given loudspeakers and room.

3.4 Sound Zone Approaches

The two main approaches in sound zone literature are pressure matching (PM) and acoustic contrast control (ACC). One of these two approaches will be used in the perceptual sound zone algorithm to be introduced in chapter 4. This section will mathematically introduce both approaches with the previously derived data model with the goal of sketching their mathematical properties. This will then later be used in section 3.5 to determine which is suitable for integration with the Par detectability discussed in chapter 2.

In the typical sound zone approach the sound zone problem is divided up into subproblems as described in section 3.2. The resulting loudspeaker input signals $x^{(l)}[n]$ are determined for a single bright-dark zone pair: the loudspeaker input signals are found such that the a target audio is achieved in the bright zone, while leakage is minimized in the dark zone. If the solution for multiple zones is desired, then multiple problems must be solved and their resulting loudspeaker input signals combined.

There is another approach however. In a multi-zone approach, the loudspeaker input signals are instead determined for jointly for all zones, rather than decomposing into bright-dark zone pairs. This is the approach that will be taken in this thesis, as it was found to be more general and simple to present. For simplicity, this thesis will limit the number of zones to two. The approach is however generalizable to any multiplicity of zones.

In a two zone multi-zone approach, the loudspeaker input signals $x^{(l)}[n]$ will be decomposed into two parts as follows:

$$x^{(l)}[n] = x_{\mathcal{A}}^{(l)}[n] + x_{\mathcal{B}}^{(l)}[n] \quad (3.6)$$

Here, $x_{\mathcal{A}}^{(l)}[n]$ and $x_{\mathcal{B}}^{(l)}[n]$ are the parts of the loudspeaker input signal responsible for reproducing the target sound pressure in zone \mathcal{A} and \mathcal{B} respectively.

Through this decomposition, it is possible to consider the sound pressure that arises due to the separate loudspeaker input signals:

$$p_{\mathcal{Z}}^{(m)}[n] = \sum_{l=0}^{N_L} \left(h^{(l,m)} * x_{\mathcal{Z}}^{(l)} \right) [n] \quad (3.7)$$

Where $\mathcal{Z} \in (\mathcal{A}, \mathcal{B})$ represents either zones. Here, $p_{\mathcal{A}}^{(m)}[n]$ and $p_{\mathcal{B}}^{(m)}[n]$ can be understood to be the pressure that arises due to playing loudspeaker input signals $x_{\mathcal{A}}^{(l)}[n]$ and $x_{\mathcal{B}}^{(l)}[n]$ respectively at a specified control point. The total sound pressure at control point m is then given by the addition of the two sound pressures:

$$p^{(m)}[n] = p_{\mathcal{A}}^{(m)}[n] + p_{\mathcal{B}}^{(m)}[n] \quad (3.8)$$

What follows is using this decomposition to describe a multi-zone variant of both a pressure matching approach in subsection 3.4.1 and acoustic contrast control approach in subsection 3.4.2.

3.4.1 Pressure Matching

The “Pressure Matching” (PM) is widely used in literature to solve the sound zone problem. In this section, a “Multi-Zone Pressure-Matching” (MZ-PM) algorithm will be derived using the previously derived data model.

In pressure matching approaches, one attempts to control the output of the loud-speaker array in such a way that resulting sound pressure in the zone matches the specified target sound pressure for that zone, while simultaneously minimizing the sound pressure that results in other zones as to minimize the interference or crosstalk between zones [1, 25].

The idea in this approach is to chose $x_{\mathcal{A}}^{(l)}[n]$ and such that the resulting pressure $p_{\mathcal{A}}^{(m)}[n]$ attains the target sound pressure $t^{(m)}[n]$ in all control points $m \in \mathcal{A}$. At the same time however, $p_{\mathcal{A}}^{(m)}[n]$ should result minimal sound pressure in all control points $m \in \mathcal{B}$. Any sound pressure resulting from $x_{\mathcal{A}}^{(l)}[n]$ in zone \mathcal{B} is can be understood as leakage, or cross-talk between zones.

Similar arguments can be given for $x_{\mathcal{B}}^{(l)}[n]$.

An optimization problem that achieves this goal can be formulated as follows:

$$\arg \min_{x_{\mathcal{A}}^{(l)}[n], x_{\mathcal{B}}^{(l)}[n] \forall l} \sum_{m \in \mathcal{A}} \left\| p_{\mathcal{A}}^{(m)}[n] - t^{(m)}[n] \right\|_2^2 + \sum_{m \in \mathcal{A}} \left\| p_{\mathcal{B}}^{(m)}[n] \right\|_2^2 + \quad (3.9)$$

$$\sum_{m \in \mathcal{B}} \left\| p_{\mathcal{B}}^{(m)}[n] - t^{(m)}[n] \right\|_2^2 + \sum_{m \in \mathcal{B}} \left\| p_{\mathcal{A}}^{(m)}[n] \right\|_2^2 \quad (3.10)$$

Here, the $\| \cdot \|_2^2$ operates denotes the squared L2-norm. In this case, this corresponds to taking the sum of the squares of the elements of a time-domain input sequence.

To further understand how optimization problem operates, consider the following definitions:

$$\text{RE}_{\mathcal{Z}} = \sum_{m \in \mathcal{Z}} \left\| p_{\mathcal{Z}}^{(m)}[n] - t^{(m)}[n] \right\|_2^2 \quad (3.11)$$

$$\text{LE}_{\mathcal{Z}} = \sum_{m \notin \mathcal{Z}} \left\| p_{\mathcal{Z}}^{(m)}[n] \right\|_2^2 \quad (3.12)$$

Here, $\text{RE}_{\mathcal{Z}}$ is the reproduction error for zone $\mathcal{Z} \in (\mathcal{A}, \mathcal{B})$. This is error corresponds to how well the achieved sound pressure $p_{\mathcal{Z}}^{(m)}[n]$ matches the target sound pressure $t^{(m)}[n]$ for all control points in the bright zone \mathcal{Z} .

$\text{LE}_{\mathcal{Z}}$ is the leakage error in zone $\mathcal{Z} \in (\mathcal{A}, \mathcal{B})$. This can be understood as the sound pressure that “leaks” into zones other than \mathcal{Z} when attempting to reproduce the target sound pressure $t^{(m)}[n]$ in zone \mathcal{Z} . This can be otherwise be considered as the “interference” or “cross-talk” between zones.

Using these new definition allows for the following rewrite of the optimization problem:

$$\arg \min_{x_{\mathcal{A}}^{(l)}[n], x_{\mathcal{B}}^{(l)}[n] \forall l} \text{RE}_{\mathcal{A}} + \text{LE}_{\mathcal{A}} + \text{RE}_{\mathcal{B}} + \text{LE}_{\mathcal{B}} \quad (3.13)$$

From this it becomes clear that this approach results in trade-off between minimizing the reproduction errors $\text{RE}_{\mathcal{Z}}$ and leakage errors $\text{LE}_{\mathcal{Z}}$.

Some pressure matching approaches attempt to control this trade-off by introducing weights for the different error terms, or by adding constraints. Choosing constraints can however be challenging as the mean square pressure error is difficult to interpret.

3.4.2 Acoustic Contrast Control

“Acoustic Contrast Control” is another widely used sound zone approach from literature. The acoustic contrast control (ACC) approach to sound zones attempts to maximize the acoustic contrast between the bright zone and the dark zone. Acoustic contrast is defined as the ratio of the total sound energy of the bright zone and the dark zone. Essentially, the goal is to maximize the difference in sound pressure level between the bright and dark zones.

In this section, a “Multi-Zone Acoustic Contrast Control” (MZ-ACC) algorithm will be described. As the previously described data model is in the time domain, this approach will take inspiration from a time-domain approach found in literature known as the broadband acoustic contrast control (BACC) approach [26, 27, 28].

In contrast to the MZ-PM approach, the MZ-ACC approach does not optimize directly over the loudspeaker input signals $x_{\mathcal{A}}^{(l)}[n]$ and $x_{\mathcal{B}}^{(l)}[n]$. Instead, it indirectly controls the loudspeaker input signals by optimizing over FIR filter coefficients $w_{\mathcal{A}}^{(l)}[n] \in \mathbb{R}^{N_w}$ and $w_{\mathcal{B}}^{(l)}[n] \in \mathbb{R}^{N_w}$. These filters are applied to the desired playback signals $s_{\mathcal{A}}^{(l)}$ and $s_{\mathcal{B}}^{(l)}$ respectively to form the final loudspeaker input signals.

This relationship between the loudspeaker input signals and the filter coefficients is thus given as follows:

$$x_{\mathcal{Z}}^{(l)}[n] = \left(w_{\mathcal{Z}}^{(l)} * s_{\mathcal{Z}} \right) [n] \quad (3.14)$$

This definition also relates the filter coefficients to the resulting sound pressure through Equation 3.4.

As mentioned, the goal of the ACC approach is to maximize the acoustic contrast between bright and dark zones, which was defined as the ratio between the sound energy in the bright and dark zones. The total sound energy in a zone will be defined as the sum of squares of the sound pressure in a control point. As such, the acoustic contrast $\text{AC}_{\mathcal{Z}}$ for a zone \mathcal{Z} can be defined as follows:

$$\text{AC}_{\mathcal{Z}} = \frac{\sum_{m \in \mathcal{Z}} \left\| p_{\mathcal{Z}}^{(m)}[n] \right\|_2^2}{\sum_{m \notin \mathcal{Z}} \left\| p_{\mathcal{Z}}^{(m)}[n] \right\|_2^2} \quad (3.15)$$

In an ACC approach the goal is to maximize the total acoustic contrast. Thus, consider the following optimization problem:

$$\arg \max_{w_{\mathcal{A}}^{(l)}[n], w_{\mathcal{B}}^{(l)}[n] \forall l} \text{AC}_{\mathcal{A}} + \text{AC}_{\mathcal{B}} \quad (3.16)$$

As mentioned, the optimization is performed over the loudspeaker filter coefficients rather than over the loudspeaker input signals.

3.5 Selection of Sound Zone Approach

In chapter 2 it was determined that the Par detectability is the perceptual model most suited for use in a perceptual sound zone algorithm. Previously, in section 3.4, two different sound zone techniques were discussed pressure matching (PM), and acoustic contrast control (ACC).

In this section, a selection will be made on what sound zone approach is most suited for integration with the Par detectability. This will be done by first reflecting on its mathematical properties in subsection 3.5.1. Next, subsection 3.5.2 will argue which approach is most promising.

3.5.1 Mathematical Properties of Par Detectability

This section will reflect on the mathematical properties of the Par detectability.

Recall from section 2.5 that the detectability $D(x[n], \varepsilon[n])$ quantifies how detectable a disturbance $\varepsilon[n] \in \mathbb{R}^{N_x}$ is in presence of a masking signal $x[n] \in \mathbb{R}^{N_x}$. The Par detectability assumes that the time-scale of its inputs are short, in the order of 20 to 200 ms.

The detectability is computed using the frequency domain representation of its input signals. To this end, $X[k]$ and $\mathcal{E}[k]$ were introduced to denote the frequency domain representations of $x[n]$ and $\varepsilon[n]$ respectively.

In subsection 2.5.4 it was found that the detectability could be expressed as in least-squares fashion as follows:

$$D(x[n], \varepsilon[n]) = \|W_x[k]\mathcal{E}[k]\|_2^2 \quad (3.17)$$

Here, perceptual weighting $W_x[k]$ models the psycho-acoustical masking effects of the masking signal $x[n]$. The weights are applied to the frequency bins of the disturbance signal $\mathcal{E}[k]$, and the sum of squares is used to determine the final detectability rating.

Note that the detectability is a convex function of the disturbance signal $\mathcal{E}[k]$ when the masking is held fixed. Convexity is a very desirable property for optimization purposes as it guarantees a unique global minimum or maximum for the optimization [23]. In addition, there are many efficient solvers for convex problems.

As such, one approach is to specify a sound zone algorithm which leverages the optimization over this disturbance signal in some way. This will be done by adopting a model for the disturbance $\mathcal{E}[k]$. For example, the disturbance could model the sound pressure error.

In summary, the detectability has the following properties that must be taken into account:

1. It is computed in the frequency domain.
2. It operates on short-time scales.

3. It is convex when optimizing over the disturbance signal $\mathcal{E}[k]$.

3.5.2 Selecting an Appropriate Sound Zone Approach

In section 3.4 two different sound zone approaches were discussed: pressure matching (PM) and acoustic contrast control (ACC). In this section it will be determined which approach of these two approaches is most promising for integration with a perceptual model.

Currently, neither approach introduced in section 3.4 operates in the frequency domain on short-time scales. However, it has been shown in literature that both ACC and PM can be adapted to meet this need.

As such, the decision between ACC and PM will depend on which approach can more naturally use the detectability. Through inspection, it was found that the PM formulation can easily be extended to include Par detectability directly.

To see this, recall in the discussion of PM given in section 3.4 that PM was shown to minimize the sum of the reproduction error in the bright zone $\text{RE}_{\mathcal{Z}}$ and the leakage to the dark zone $\text{LE}_{\mathcal{Z}}$. Their definitions were given as follows:

$$\text{RE}_{\mathcal{Z}} = \sum_{m \in \mathcal{Z}} \left\| p_{\mathcal{Z}}^{(m)}[n] - t^{(m)}[n] \right\|_2^2 \quad (3.18)$$

$$\text{LE}_{\mathcal{Z}} = \sum_{m \notin \mathcal{Z}} \left\| p_{\mathcal{Z}}^{(m)}[n] \right\|_2^2 \quad (3.19)$$

Here, the reproduction error $\text{RE}_{\mathcal{Z}}$ can be understood as the total error between the achieved and target sound pressure in the bright zone of zone \mathcal{Z} . The leakage $\text{LE}_{\mathcal{Z}}$ can be understood as the sound pressure that arises in zones other than \mathcal{Z} , essentially “leaking” into other zones, when reproducing the target in the bright zone.

Both can be thus be considered as “errors” that we wish to minimize. Viewing it this way, one choice for model these errors as the disturbance signals $\varepsilon[n]$. To this end, consider the perceptual reproduction error $\text{PRE}_{\mathcal{Z}}$ and the perceptual leakage error $\text{PLE}_{\mathcal{Z}}$. From now, take no mind to the fact that the quantities defined so far do not satisfy the short-time requirements. The perceptual errors are defined as:

$$\text{PRE}_{\mathcal{Z}} = \sum_{m \in \mathcal{Z}} D(t^{(m)}[n], p_{\mathcal{Z}}^{(m)}[n] - t^{(m)}[n]) \quad (3.20)$$

$$\text{PLE}_{\mathcal{Z}} = \sum_{m \notin \mathcal{Z}} D(t^{(m)}[n], p_{\mathcal{Z}}^{(m)}[n]) \quad (3.21)$$

The perceptual reproduction error $\text{PRE}_{\mathcal{Z}}$ is determined by the detectability of the sound pressure difference between the achieved sound pressure $p_{\mathcal{Z}}^{(m)}[n]$ and the target sound pressure $t^{(m)}[n]$ for the control point m . This sound pressure difference can be understood as the “error” in sound pressure.

The target sound pressure is used as the masking signal, in doing so the masking properties of the target sound pressure is used. The $\text{PRE}_{\mathcal{Z}}$ models how detectable the deviation from target sound pressure is. One interpretation of minimizing $\text{PRE}_{\mathcal{Z}}$ is finding the sound pressure difference which is least detectable.

Similarly, the perceptual leakage error $\text{PLE}_{\mathcal{Z}}$ is determined by how detectable the sound pressure of zone \mathcal{Z} is in other zones. Note as $m \notin \mathcal{Z}$, this definition considers all points outside of the bright zone for \mathcal{Z} . Here, the masking signal is again chosen as the target sound pressure of the control point in question. $\text{PLE}_{\mathcal{Z}}$ models how detectable the leakage of zone \mathcal{Z} is in presence of the intended target sound pressure. Minimizing $\text{PLE}_{\mathcal{Z}}$ could thus be understood as finding the leakage that is minimally-detectable in presence of the target sound pressure.

These perceptually modified errors were found to be a promising way of combining pressure matching and the Par detectability, due to the mathematical similarities between the definition of the Par detectability given by Equation 3.17 and the errors used in the pressure matching problem given by Equation 3.18 and Equation 3.19.

In the ACC approach, it was found that there was not an as obvious choice for the disturbance and masking signals. As such, the PM approach will be used in the perceptual sound zone algorithm in chapter 4. As mentioned, the PM approaches given by section 3.4 are not given in neither the frequency domain nor operate on short time scales. This will be addressed in the same chapter, in section 4.2 and section 4.3.

3.6 Conclusion

Implementation of Perceptual Sound Zone Algorithms

4

4.1 Introduction

This chapter seeks to combine the selected perceptual model and the selected sound zone algorithm. In chapter 2, it was found that the Par detectability was best suited. Based on the constraints posed by this perceptual model, chapter 3 found that a pressure matching approach was best suited, and gave an implementation that can be used as a foundation to build the perceptual algorithm on.

This chapter is structured as follows.

4.2 Short-Time Reformulation of Pressure Matching

In section 3.5 it was determined that the pressure matching approach introduced in ?? was suitable for integration with the Par detectability. The Par detectability assumes that its inputs are given in short-time segments of 20 to 200 milliseconds. The previously discussed pressure matching approach however attempts to compute the loudspeaker input signals in their entirety.

Thus, in this section, the existing pressure matching approach is reformulated to operate on short-time segments through a “block-based” approach, where small “blocks” of samples are processed block-by-block rather than the entirety of the signal.

Aside from ease of integrating a perceptual model, there is another motivation for using a block-based approach. In the preceding section it is assumed that the desired playback signals $s_{\mathcal{A}}[n]$ and $s_{\mathcal{B}}[n]$ were known in their entirety. In practice however, this is not a valid assumption as a user can change the desired playback content in real-time. This is the case for example when a user changes the song they are playing on their system.

Instead, a block-based approach can be used where the incoming samples of the desired playback signals are used in real-time as they become available. The system buffers a block of H incoming samples, and then solves the sound zone problem finding the newest block of loudspeaker input signals.

For these reasons, this section will adapt existing pressure matching approach introduced in section 3.4 to accommodate for block based processing. This will be done by first modeling the knowledge limitation of the block-based processing on the desired playback signals.

4.2.1 Modeling Block Based Knowledge Limitations

For the block-processing based sound zone approach, the incoming samples of the desired playback signals $s_{\mathcal{Z}}[m]$ for $\mathcal{Z} \in \{\mathcal{A}, \mathcal{B}\}$ are assumed to be buffered into blocks of size H . This means that the system waits until it has a block of size H before revealing the new samples to the sound zone system.

This introduces a knowledge limitation in the desired playback signals, which requires the sound zone algorithm to be redesigned. In this section, this knowledge limitation will be modeled.

The sound zone system only has knowledge of all previous blocks and the most recent block, indexed by block index μ . The relation between the global time index n and block index μ is given as follows:

$$\mu(n) = \lfloor n/H \rfloor \quad (4.1)$$

Thus at a time n , up to and including the μ^{th} blocks of desired playback signals $s_{\mathcal{Z}}[n]$ are known.

Let the available knowledge of the desired playback signal with knowledge of blocks up to μ be denoted as $\tilde{s}_{\mathcal{Z}}^R[n, \mu]$. It can be defined as follows:

$$\tilde{s}_{\mathcal{Z}}^R[n, \mu] = \sum_{k=-\infty}^{\mu} \tilde{s}_{\mathcal{Z},k}^R[n] w_H^R[n - kH] \quad (4.2)$$

Here, $w_H^R[n] \in \mathbb{R}^H$ is a non-causal rectangular window with support $-H + 1 \leq n \leq 0$. Furthermore, $\tilde{s}_{\mathcal{Z},k}^R[n]$ is the k^{th} block of the playback signal $s_{\mathcal{Z}}[n]$. As such, $\tilde{s}_{\mathcal{Z},k}^R[n] = s_{\mathcal{Z}}[n]$ for $-H + 1 + kH \leq n \leq kH$ (and zero elsewhere).

One way of interpreting the equation above is as a projection of $s_{\mathcal{Z}}[n]$ on a basis spanned by shifted non-overlapping rectangular windows $w_H^R[n - kH]$. Here, $\tilde{s}_{\mathcal{Z},k}^R[n]$ can be thought of as the coefficients for the basis functions resulting from the projection.

Expressing it this way reveals that a more general approach is possible: as shown, desired playback signals are projected onto non-causal non-overlapping rectangular windows of size H . This model can however be generalized to a larger class of windows.

Let $w[n] \in \mathbb{R}^{N_w}$ denote a non-causal window with support $-N_w + 1 \leq n \leq 0$. Here, $w[n]$ is chosen such that it is COLA-condition compliant for hop size H . The COLA condition requires that the sum of all shifted windows add to unity for all samples n for a given hop size H . It is given as follows:

$$\sum_{k=-\infty}^{\infty} w[n - kH] = 1 \quad \forall n \quad (4.3)$$

Projecting $s_{\mathcal{Z}}[n]$ onto a basis spanned by the shifted windows $w[n - kH]$ results in the following:

$$s_{\mathcal{Z}}[n] = \sum_{k=-\infty}^{\infty} s_{\mathcal{Z}}[n] w[n - kH] \quad (4.4)$$

$$= \sum_{k=-\infty}^{\infty} \tilde{s}_{\mathcal{Z},k}[n] w[n - kH] \quad (4.5)$$

Here, $\tilde{s}_{\mathcal{Z},k}[n] = s_{\mathcal{Z}}[n]$ for $-N_w + 1 + kH \leq n \leq kH$ and zero elsewhere. The windows decimate the signal $s_{\mathcal{Z}}[n]$ into segments of size N_w . Due to the COLA condition, this the segments reconstruct the

This allows us to express the desired playback signal with knowledge up to block μ as follows:

$$\tilde{s}_{\mathcal{Z}}[n, \mu] = \sum_{k=-\infty}^{\mu} \tilde{s}_{\mathcal{Z},k}[n] w[n - kH] \quad (4.6)$$

This form will converge to the real desired playback signal as $\mu \rightarrow \infty$. Aside from being more general, this approach allows for use of overlapping windows. This was found to be beneficial, as overlapping windows were found to reduce edge effects in the resulting loudspeaker input signals.

4.2.2 Block Based Loudspeaker Input Signal Computation

As mentioned in the introduction, the block-based approach aims to compute the loudspeaker input signals at the same rate as the desired playback signals $s_{\mathcal{Z}}[n]$ are received. That is to say: when a new block μ of the desired playback signal is revealed, a new block μ of loudspeaker input signals need be computed.

In order to compute the loudspeaker input signals in this block-wise fashion, a similar decomposition as for the desired playback signals is performed. Consider decomposing the loudspeaker input signals in blocks through the windows as follows:

$$\tilde{x}_{\mathcal{Z}}^{(l)}[n, \mu] = \sum_{k=-\infty}^{\mu} \tilde{x}_{\mathcal{Z},k}^{(l)}[n]w[n - kH] \quad (4.7)$$

$$= \sum_{k=-\infty}^{\mu-1} \tilde{x}_{\mathcal{Z},k}^{(l)}[n]w[n - kH] + \tilde{x}_{\mathcal{Z},\mu}^{(l)}[n]w[n - \mu H] \quad (4.8)$$

$$= \tilde{x}_{\mathcal{Z}}^{(l)}[n, \mu - 1] + \tilde{x}_{\mathcal{Z},\mu}^{(l)}[n]w[n - \mu H] \quad (4.9)$$

The equation above shows that $\tilde{x}_{\mathcal{Z}}^{(l)}[n]$ can be computed by recursively computing its coefficients $\tilde{x}_{\mathcal{Z},k}^{(l)}$ and adding them. Note that this recursive property also holds for the desired playback signal $\tilde{s}_{\mathcal{Z}}[n]$.

The approach to find $\tilde{x}_{\mathcal{Z}}^{(l)}[n]$ is now as follows. When the μ^{th} block of the desired playback signals is revealed, this will allow for the calculation of the corresponding target sound pressure. Next, the μ^{th} coefficient $\tilde{x}_{\mathcal{Z},\mu}^{(l)}$ is computed such that $\tilde{x}_{\mathcal{Z}}^{(l)}[n]$ attains this target sound pressure.

How $\tilde{x}_{\mathcal{Z}}^{(l)}[n]$ and $\tilde{s}_{\mathcal{Z}}[n]$ relate to the resulting and target sound pressure respectively is the topic of the next section.

4.2.3 Block Based Sound Pressure

As discussed previously, the Multi-Zone Pressure-Matching (MZ-PM) algorithm attempts to control the loudspeaker input signals $x_{\mathcal{Z}}^{(l)}[n]$ for $\mathcal{Z} \in \{\mathcal{A}, \mathcal{B}\}$ such that the resulting sound pressure matches a specified target sound pressure $t^{(m)}[n]$ at all control points m .

Previous sections gave block-based versions of the loudspeaker input signals and the desired playback signals. In this section, block-based versions of the resulting and target sound pressure will be given.

The block-based target sound pressure will be denoted by $\tilde{t}^{(m)}[n, \mu]$, and can be defined by simply substituting the definition for the block-based desired playback

signal $\tilde{s}_{\mathcal{Z}}[n, \mu]$ into the definition of the target pressure Figure A.3:

$$\tilde{t}^{(m)}[n, \mu] = \sum_{l=0}^{N_L} \left(h^{(l,m)} * \tilde{s}_{\mathcal{Z}} \right) [n] \quad (4.10)$$

$$= \sum_{l=0}^{N_L} \sum_{k=-\infty}^{\mu} \sum_{m=0}^{N_h-1} h^{(l,m)}[m] \tilde{s}_{\mathcal{Z},k}[n-m] w[n-kH-m] \quad (4.11)$$

$$= \tilde{t}^{(m)}[n, \mu-1] + \sum_{l=0}^{N_L} \sum_{m=0}^{N_h-1} h^{(l,m)}[m] \tilde{s}_{\mathcal{Z},\mu}[n-m] w[n-\mu H-m] \quad (4.12)$$

The definition above holds for all points $m \in \mathcal{Z}$, the points contained in zone \mathcal{Z} . As can be seen, the block based target sound pressure for the block μ can be computed recursively by adding the contribution of the newest block $\tilde{s}_{\mathcal{Z},\mu}[n]$ the target sound pressure of the previous block.

The block-based resulting sound pressure will be denoted by $\tilde{p}_{\mathcal{Z}}^{(m)}[n, \mu]$, and can be defined by simply substituting the definition for the block-based loudspeaker input signals $\tilde{x}_{\mathcal{Z}}^{(l)}[n, \mu]$ into the definition of the resulting pressure Figure A.3. This results in the following:

$$\tilde{p}_{\mathcal{Z}}^{(m)}[n, \mu] = \sum_{l=0}^{N_L} \left(h^{(l,m)} * \tilde{x}_{\mathcal{Z}}^{(l)} \right) [n] \quad (4.13)$$

$$= \sum_{l=0}^{N_L} \sum_{k=-\infty}^{\mu} \sum_{m=0}^{N_h-1} h^{(l,m)}[m] \tilde{x}_{\mathcal{Z},k}^{(l)}[n-m] w[n-kH-m] \quad (4.14)$$

$$= \tilde{p}_{\mathcal{Z}}^{(m)}[n, \mu-1] + \sum_{l=0}^{N_L} \sum_{m=0}^{N_h-1} h^{(l,m)}[m] \tilde{x}_{\mathcal{Z},\mu}^{(l)}[n-m] w[n-\mu H-m] \quad (4.15)$$

The definition above again holds for all points $m \in \mathcal{Z}$. As can be seen, the block based resulting sound pressure for the block μ can also be computed recursively.

With this, all quantities required for the block-based formulation of the Multi-Zone Pressure-Matching (MZ-PM) approach are known. The next section will use these quantities to state the block-based MZ-PM algorithm.

4.2.4 Derivation of Block-Based Multi-Zone Pressure-Matching

After translating the loudspeaker input signals and the target sound pressure into their block-wise counterparts, the Block-Based Multi-Zone Pressure-Matching (BB-MZ-PM) algorithm can be stated.

As mentioned before, the approach that will be taken is to compute the μ^{th} coefficient of the loudspeaker input signal $\tilde{x}_{\mathcal{Z},\mu}^{(l)}$ such that the resulting sound pressure $\tilde{p}_{\mathcal{Z}}^{(m)}[n, \mu]$ best matches the target sound pressure $\tilde{t}^{(m)}[n, \mu]$.

Note that in this approach only the most recent loudspeaker coefficients $\tilde{x}_{\mathcal{Z},\mu}^{(l)}$ are being controlled. The previous coefficients are assumed to have already been played. As such, they are held fixed.

The block-based optimization problem can be found by simply replacing all quantities in the previously derived optimization problem with their block-based counterparts. The problem is given as follows:

$$\arg \min_{\tilde{x}_{\mathcal{A},\mu}^{(l)}, \tilde{x}_{\mathcal{B},\mu}^{(l)} \forall l} \sum_{m \in \mathcal{A}} \left\| \tilde{p}_{\mathcal{A}}^{(m)}[n, \mu] - \tilde{t}_{\mu}^{(m)}[n, \mu] \right\|_2^2 + \sum_{m \in \mathcal{A}} \left\| \tilde{p}_{\mathcal{B}}^{(m)}[n, \mu] \right\|_2^2 + \quad (4.16)$$

$$\sum_{m \in \mathcal{B}} \left\| \tilde{p}_{\mathcal{B}}^{(m)}[n, \mu] - \tilde{t}_{\mu}^{(m)}[n, \mu] \right\|_2^2 + \sum_{m \in \mathcal{B}} \left\| \tilde{p}_{\mathcal{A}}^{(m)}[n, \mu] \right\|_2^2 \quad (4.17)$$

Note that this problem implicitly contains the target sound pressure and resulting sound pressure of the previous blocks $\mu - 1$ due to the aforementioned recursive definitions. As a result, the history of what has been transmitted by the loudspeaker previously is included in the optimization.

The problem above is solved recursively for all loudspeaker input signal coefficients $\tilde{x}_{\mathcal{A},\mu}^{(l)}[n]$ and $\tilde{x}_{\mathcal{B},\mu}^{(l)}[n]$ as new blocks $\tilde{s}_{\mathcal{A},\mu}[n]$ and $\tilde{s}_{\mathcal{B},\mu}[n]$ are revealed. The final loudspeaker input signals $\tilde{x}_{\mathcal{Z}}^{(l)}[n, \infty]$ can then be found by means of Figure A.3.

4.3 Frequency Domain Reformulation of Pressure Matching

In the previous section, a block-based multi-zone pressure matching algorithm was derived. When deriving this algorithm it was stated that its advantages are twofold.

Firstly, one advantage of using this algorithm over its non block-based counterpart is that it can work in real-time. Secondly, the block-based approach can operate on short time-scales. This is useful, as the perceptual model that we wish to integrate operates on time-scales of the order of 20 to 200 milliseconds.

There is however an additional adjustment that needs to be made before the perceptual model can be integrated. Currently, the block based pressure matching algorithm minimizes the sound pressure errors in the time-domain, whereas the detectability is computed in the frequency domain. In order to integrate the perceptual model and the sound zone algorithm, they must operate in the same domain.

In this case, it was chosen to use convert the existing sound zone algorithm to one that operates in the frequency domain. This was chosen as to remain close to the definition of the detectability. Approximating detectability in the time domain was not explored further and could potentially be of interest for future work.

To this end, this section will adjust the existing time domain block based pressure matching algorithm to an equivalent frequency domain formulation. This will be done by first introducing a suitable transformation relating the time and frequency domain quantities. The transformed quantities can then be used to define the frequency domain version of the block based pressure matching algorithm.

4.3.1 Frequency Domain Transformation

A suitable transform to obtain the frequency domain representation of a signal is the DFT. However, it is important to take a number of precautions before applying the DFT directly, as the computation of the sound pressures used in the optimization problem introduced previously involves taking the linear convolution of the loudspeaker input signals with the room impulse responses.

Time domain circular convolution can be computed in the frequency domain through the Hadamard product. Time domain circular convolution coincides with time domain linear convolution only if the two operands are zero-padded sufficiently. To be specific, both operands need be zero-padded to the length of the resulting convolution.

As such, the frequency domain transform requires this zero padding to be built in. The convolutions described in the previous chapter are between the window coefficients of size N_w and the room impulse responses of size N_h . Thus, the both must be zero padded to convolution length $N_w + N_h - 1$ before going to the frequency domain.

A suitable transform is thus given by a $N_w + N_h - 1$ point DFT.

$$X[k] = \sum_{n=0}^{N_w+N_h-2} x[n] \exp\left(\frac{-j2\pi kn}{N_w + N_h - 1}\right) \quad (4.18)$$

Here, $x[n]$ and $X[k]$ are the time- and frequency-domain representations of an arbitrary sequence.

4.3.2 Quantities the Frequency Domain

In this section, the quantities used in the Block-Based Multi-Zone Pressure-Matching approach will be converted to their frequency domain counterparts.

Essentially, this involves converting the sound pressures $\tilde{p}_{\mathcal{Z}}^{(m)}[n, \mu]$ and $\tilde{t}^{(m)}[n, \mu]$ to their frequency domain versions given by $\tilde{P}_{\mathcal{Z},\mu}^{(m)}[k]$ and $\tilde{T}_{\mu}^{(m)}[k]$ respectively.

This can be done by applying the previously derived transform directly to these quantities. This results in the following expressions.

$$\tilde{T}^{(m)}[k, \mu] = \tilde{T}^{(m)}[k, \mu - 1] + \sum_{l=0}^{N_L} H^{(l,m)}[k] \tilde{S}_{\mathcal{Z},\mu}[k] \quad (4.19)$$

$$\tilde{P}_{\mathcal{Z}}^{(m)}[k, \mu] = \tilde{P}_{\mathcal{Z}}^{(m)}[k, \mu - 1] + \sum_{l=0}^{N_L} H^{(l,m)}[k] \tilde{X}_{\mathcal{Z},\mu}^{(l)}[k] \quad (4.20)$$

Here, $H^{(l,m)}[k] \in \mathbb{C}^{N_w+N_h-1}$ is the transformed version of the room impulse responses. Furthermore, $\tilde{S}_{\mathcal{Z},\mu}[k] \in \mathbb{C}^{N_w+N_h-1}$ and $\tilde{X}_{\mathcal{Z},\mu}^{(l)}[k] \in \mathbb{C}^{N_w+N_h-1}$ are the frequency domain versions of the desired playback signal and the loudspeaker input signal, which are defined as follows:

$$\tilde{S}_{\mathcal{Z},\mu}[k] = \sum_{n=0}^{N_w+N_h-2} \tilde{s}_{\mathcal{Z},\mu}[n] w[n - \mu H] \exp\left(\frac{-j2\pi kn}{N_w + N_h - 1}\right) \quad (4.21)$$

$$\tilde{X}_{\mathcal{Z},\mu}^{(l)}[k] = \sum_{n=0}^{N_w+N_h-2} \tilde{x}_{\mathcal{Z},\mu}^{(l)}[n] w[n - \mu H] \exp\left(\frac{-j2\pi kn}{N_w + N_h - 1}\right) \quad (4.22)$$

Note that the window is implicitly included in the transformed quantities. This is done for ease of notation.

4.3.3 Proposed Frequency Domain Approach

Using the previously derived quantities, it is possible express the frequency domain version of the Block-Based Multi-Zone Pressure-Matching (BB-MZ-PM) approach as follows:

$$\arg \min_{\tilde{x}_{\mathcal{A},\mu}^{(l)}, \tilde{x}_{\mathcal{B},\mu}^{(l)} \forall l} \sum_{m \in \mathcal{A}} \left\| \tilde{P}_{\mathcal{A}}^{(m)}[k, \mu] - \tilde{T}^{(m)}[k, \mu] \right\|_2^2 + \sum_{m \in \mathcal{A}} \left\| \tilde{P}_{\mathcal{B}}^{(m)}[k, \mu] \right\|_2^2 + \quad (4.23)$$

$$\sum_{m \in \mathcal{B}} \left\| \tilde{P}_{\mathcal{B}}^{(m)}[k, \mu] - \tilde{T}^{(m)}[k, \mu] \right\|_2^2 + \sum_{m \in \mathcal{B}} \left\| \tilde{P}_{\mathcal{A}}^{(m)}[k, \mu] \right\|_2^2 \quad (4.24)$$

Note how the optimization is still performed over the time domain signal. This was done to constrain the loudspeaker input signal coefficient to size N_w , as that is an assumption made by the frame-based processing.

In principal, this introduces more complexity than solving directly over the frequency domain loudspeaker input coefficient $\tilde{X}_{\mathcal{Z},\mu}^{(l)}[k]$. This however introduces issues as it requires the truncation of the time-domain version to N_w samples, which was found to introduce artifacts.

4.4 Unconstrained Perceptual Pressure Matching

In chapter 2 the Par detectability $D(x[n], \varepsilon[n])$ was introduced as the most promising perceptual model. Later, in section 3.5 it was found that a pressure matching could be formulated using said perceptual model. This section will detail the combination of the two, resulting in a detectability-based pressure matching algorithm.

4.4.1 Summary of Par Detectability

The Par detectability $D(x[n], \varepsilon[n])$ quantifies how detectable a disturbance signal $\varepsilon[n]$ is in presence of masking signal $x[n]$. Essentially, it assigns a grade to how well a human can detect $\varepsilon[n]$ when also listening to $x[n]$.

It was shown in section 2.5 that the Par detectability $D(x[n], \varepsilon[n])$ can be expressed as follows:

$$D(x[n], \varepsilon[n]) = \|W_x[k]\mathcal{E}[k]\|_2^2 \quad (4.25)$$

Here, perceptual weighting $W_x[k]$ defines a frequency weighting dependant on the masking properties of the masking signal $x[n]$. The perceptual weighting is applied to the frequency domain representation of the disturbance signal $\mathcal{E}[k]$ to obtain the detectability grade.

If the masking signal $x[n]$ is held constant, the detectability is convex in the disturbance signal $\varepsilon[n]$.

4.4.2 Summary of Pressure Matching

In pressure matching, the goal is to minimize the sound pressure leaking into the dark zone while also minimizing the difference between the target and achieved sound pressure in the bright zone. In section 3.4 it was shown that the pressure matching algorithm can be given as follows:

$$\arg \min_{x_{\mathcal{A}}^{(l)}[n], x_{\mathcal{B}}^{(l)}[n] \forall l} \text{RE}_{\mathcal{A}} + \text{LE}_{\mathcal{A}} + \text{RE}_{\mathcal{B}} + \text{LE}_{\mathcal{B}} \quad (4.26)$$

In this optimization problem, the loudspeaker input signals $x_{\mathcal{A}}^{(l)}[n]$ and $x_{\mathcal{B}}^{(l)}[n]$ for all loudspeakers l are determined such that they minimize the reproduction errors $\text{RE}_{\mathcal{A}}$ and $\text{RE}_{\mathcal{B}}$, and the leakage errors $\text{LE}_{\mathcal{A}}$ and $\text{LE}_{\mathcal{B}}$.

The reproduction errors $\text{RE}_{\mathcal{A}}$ and $\text{RE}_{\mathcal{B}}$ represent the error between target and achieved sound pressure, whereas the leakage errors $\text{LE}_{\mathcal{A}}$ and $\text{LE}_{\mathcal{B}}$ represent the leakage of bright zone content to the dark zone. Recall that they were defined as follows:

$$\text{RE}_{\mathcal{Z}} = \sum_{m \in \mathcal{Z}} \left\| p_{\mathcal{Z}}^{(m)}[n] - t^{(m)}[n] \right\|_2^2 \quad (4.27)$$

$$\text{LE}_{\mathcal{Z}} = \sum_{m \notin \mathcal{Z}} \left\| p_{\mathcal{Z}}^{(m)}[n] \right\|_2^2 \quad (4.28)$$

Here, $p_Z^{(m)}[n]$ is the achieved sound pressure at control point m and $t^{(m)}[n]$ is the target sound pressure for control point m .

As the detectability operates on short-time scales, a short-time version of pressure matching was proposed in

4.4.3 Detectability Based Pressure Matching

Previously, section 3.5 found that pressure matching approach could be formulated using the detectability. Consider the following definitions for the perceptual reproduction error and the perceptual leakage error:

$$\text{PRE}_Z = \sum_{m \in Z} D(t^{(m)}[n], p_Z^{(m)}[n] - t^{(m)}[n]) \quad (4.29)$$

$$\text{PLE}_Z = \sum_{m \notin Z} D(t^{(m)}[n], p_Z^{(m)}[n]) \quad (4.30)$$

Rather than optimizing over the

4.4.4 Proposed Approach

Essentially, the approach is to replace the norms in the equation above by the perceptually weighted equivalent. For each norm, we consider the masking effects determined by the target sound pressure for the respective point m . This results in the following cost function:

$$\arg \min_{\tilde{x}_{\mathcal{A},\mu}^{(l)}[n], \tilde{x}_{\mathcal{B},\mu}^{(l)}[n] \forall l} \sum_{m \in \mathcal{A}} \left\| \mathbf{w} \left(\tilde{T}^{(m)}[k, \mu] \right) \left[\tilde{P}_{\mathcal{A}}^{(m)}[k, \mu] - \tilde{T}^{(m)}[k, \mu] \right] \right\|_2^2 + \quad (4.31)$$

$$\sum_{m \in \mathcal{A}} \left\| \mathbf{w} \left(\tilde{T}^{(m)}[k, \mu] \right) \left[\tilde{P}_{\mathcal{B}}^{(m)}[k, \mu] \right] \right\|_2^2 + \quad (4.32)$$

$$\sum_{m \in \mathcal{B}} \left\| \mathbf{w} \left(\tilde{T}^{(m)}[k, \mu] \right) \left[\tilde{P}_{\mathcal{B}}^{(m)}[k, \mu] - \tilde{T}^{(m)}[k, \mu] \right] \right\|_2^2 + \quad (4.33)$$

$$\sum_{m \in \mathcal{B}} \left\| \mathbf{w} \left(\tilde{T}^{(m)}[k, \mu] \right) \left[\tilde{P}_{\mathcal{A}}^{(m)}[k, \mu] \right] \right\|_2^2 \quad (4.34)$$

The cost function behaves the same way, except all terms are now psycho-acoustically weighted. The weighting vectors are chosen based on the target sound pressure in the relevant control point m . This essentially exploits the psycho-acoustical masking properties of the target signal.

4.5 Constrained Perceptual Pressure Matching

It turns out that the perceptual minimization approach doesn't always trade-off nicely between interference suppression and reproduction error minimization. In order to do so, we can move certain terms to the constraints. That is done below:

$$\arg \min_{\tilde{x}_{\mathcal{A},\mu}^{(l)}[n], \tilde{x}_{\mathcal{B},\mu}^{(l)}[n] \forall l} \sum_{m \in \mathcal{B}} \left\| \mathbf{w} \left(\tilde{T}^{(m)}[k, \mu] \right) \left[\tilde{P}_{\mathcal{A}}^{(m)}[k, \mu] \right] \right\|_2^2 + \quad (4.35)$$

$$\sum_{m \in \mathcal{A}} \left\| \mathbf{w} \left(\tilde{T}^{(m)}[k, \mu] \right) \left[\tilde{P}_{\mathcal{B}}^{(m)}[k, \mu] \right] \right\|_2^2 \quad (4.36)$$

$$\text{subject to} \quad \left\| \mathbf{w} \left(\tilde{T}^{(m)}[k, \mu] \right) \left[\tilde{P}_{\mathcal{A}}^{(m)}[k, \mu] - \tilde{T}^{(m)}[k, \mu] \right] \right\|_2^2 \leq Q \quad \forall m \in \mathcal{B} \quad (4.37)$$

$$\left\| \mathbf{w} \left(\tilde{T}^{(m)}[k, \mu] \right) \left[\tilde{P}_{\mathcal{B}}^{(m)}[k, \mu] - \tilde{T}^{(m)}[k, \mu] \right] \right\|_2^2 \leq Q \quad \forall m \in \mathcal{A} \quad (4.38)$$

Here, we are minimizing the perceptually-weighted leakage, subject to constraints limiting the detectability of the reproduction error per point m .

4.6 Conclusion

Evaluation of Perceptual Sound Zone Algorithms

5

5.1 Introduction




Figure 5.1: The room setup used in the simulations for the evaluation of the algorithms.

5.2 Evaluation Methodology

In chapter 3 and chapter 4 various sound zone algorithms were defined. First, in chapter 3 a non-perceptual pressure matching approach was introduced. In the following chapter two perceptual sound zone algorithms were introduced, namely unconstrained perceptual pressure matching and perceptually constrained pressure matching.

This section will discuss the general approach for evaluation of the previously derived algorithms. All algorithms will be evaluated by means of simulation. As all derived algorithms were found to be computationally intensive, many of the design consideration for the simulations were chosen such to keep the computational complexity low.

5.2.1 Room Setup

All sound zone algorithms were evaluated in a simulated square room of 5 by 5 meters, with a ceiling height of 3.4 meters. The room contained 4 loudspeakers placed in the corners of the room at a height of 1.2 meters.

Two zones each consisting of two control points are in the middle of the room. In order to obtain a challenging problem, the zones are placed in close proximity, which the two closest points 0.5 meters apart.

An image depicting the entire setup is given in Figure 5.1.

5.2.2 Content Selection

For the evaluation speech audio content was used. The motivation for this is that it is computationally intensive to run the algorithms at higher sampling rate such as 48 kHz, and speech signals can be represented well at lower sampling rates. In addition to this, the objective speech quality measures described in chapter 2 were found to be more robust. Many general audio quality measures were found to have little to no official openly available implementations.

A set of 4 speech signals is used for evaluation. For each experiment, a speech

signal was assigned to each zone. All possible combinations of the speech signals were made, resulting in 12 total experiments.

5.2.3 Simulation Outputs

After running the simulations, all output will be written to file. For ease of evaluation, the following sound pressures is available:

- **Target Sound Pressure per Control Point m :**
This is the desired sound pressure per control point m . This can be understood as the sound pressure that the algorithm attempts to attain. In the derivations of chapter 3, this was denoted by $t^{(m)}[n]$.
- **Achieved Target Sound Pressure per Control Point m :**
This is the target sound pressure achieved by the algorithm per control point m , excluding the interference coming from other zones. This quantity allows for the evaluation of the quality of intended sound pressure per control point, isolated from all interference. In the two zone case for any $m \in A$, this corresponds to $p_A^{(m)}[n]$ in previous derivations.
- **Achieved Interfering Sound Pressure per Control Point m :**
This is the interference sound pressure achieved by the algorithm per control point m , isolated the target sound pressure that was intended for that control point. This allows for analysis of the interference in isolation from the intended sound pressure. Again for the two zone case for any $m \in A$, this corresponds to $p_B^{(m)}[n]$.
- **Total Achieved Sound Pressure per Control Point m :**
This is the total sound pressure per control point m achieved by the algorithm. This can be understood as the sound pressure that would be experienced by a listener standing in the position of point m , and is equal to the sum over the achieved target sound pressure and the achieved interfering sound pressure previously discussed. In the derivations for the two zone case, this corresponds to $p_A^{(m)}[n] + p_B^{(m)}[n]$ for any m .

The achieved target sound pressure and the achieved interfering sound pressure can thought of as decompositions of the total achieved sound pressure. This decomposition was found to be useful during evaluation.

5.2.4 Evaluation Criteria

In section 2.3 a literature review into various perceptual grading models from literature was discussed. As mentioned before, speech audio is used as the content for the sound zones. As such, it makes sense to use metrics that specialize in speech for evaluation.

For this reason, one of the metrics that will be used is the Perceptual Evaluation of Speech Quality (PESQ). As described in subsection 2.3.1, PESQ is a metric which grades the quality of a degraded speech signal with respect to a reference speech

signal. The resulting quality grade will be between 0 and 5, where 5 is the highest obtainable grade.

PESQ will be used to evaluate two aspects of the result.

- **PESQ of the Achieved Target Sound Pressure with respect to the Target Sound Pressure per Control Point m :**

As described, the achieved target sound pressure is the sound pressure per control point m , without the interference due to the other zones.

This will henceforth referred to as the “**Target PESQ**”. This quantity thus describes the quality of the intended sound pressure in isolation from interference.

- **PESQ of the Total Achieved Sound Pressure with respect to the Target Sound Pressure per Control Point m :**

As previously described, the total achieved sound pressure is the sound pressure that would be audible to an individual present at control point m . This includes the intended sound pressure, but also the interference coming from other zones.

This quantity will be referred to as the “**Total PESQ**” and represents the quality of the total experience per control point m .

Another metric that will be used for evaluation is the Distraction model also introduced in subsection 2.3.1. This model grades how distracting an interferer is in presence of target audio. The grade uses a scale from 0 to 100, where 100 is considered maximally distracting.

The distraction will be used as follows:

- **Distraction of the Achieved Interfering Sound Pressure with respect to the Achieved Target Sound Pressure per Control Point m :**

This quantifies how distracting the interfering sound pressure is when listening to the intended sound pressure per control point m . This is a way of quantifying how disturbing the interference of the resulting algorithms is.

This will simply be referred to as the “**Distraction**” per control point from now on.

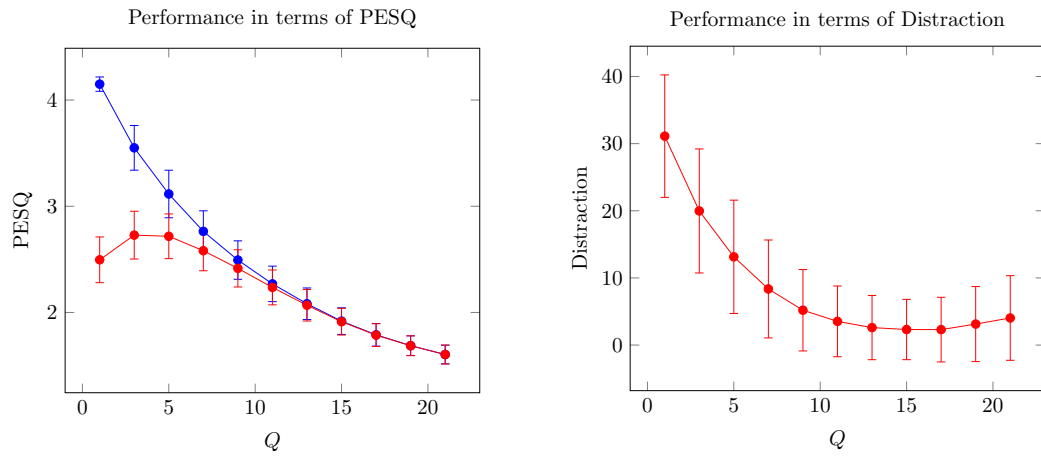


Figure 5.2

5.3 Evaluation of Proposed Algorithms

5.4 Conclusion

Bibliography

- [1] T. Betlehem, W. Zhang, M. A. Poletti, and T. D. Abhayapala, “Personal sound zones: Delivering interface-free audio to multiple listeners,” *IEEE Signal Processing Magazine*, vol. 32, no. 2, pp. 81–91, 2015.
- [2] J. Donley and C. H. Ritz, “Multizone reproduction of speech soundfields: A perceptually weighted approach,” 2015.
- [3] T. Lee, J. K. Nielsen, and M. G. Christensen, “Towards perceptually optimized sound zones: A proof-of-concept study,” in *ICASSP 2019-2019 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2019, pp. 136–140.
- [4] —, “Signal-adaptive and perceptually optimized sound zones with variable span trade-off filters,” *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, vol. 28, pp. 2412–2426, 2020.
- [5] C. H. Taal, R. C. Hendriks, R. Heusdens, and J. Jensen, “An algorithm for intelligibility prediction of time–frequency weighted noisy speech,” *IEEE Transactions on Audio, Speech, and Language Processing*, vol. 19, no. 7, pp. 2125–2136, 2011.
- [6] A. W. Rix, J. G. Beerends, M. P. Hollier, and A. P. Hekstra, “Perceptual evaluation of speech quality (pesq)-a new method for speech quality assessment of telephone networks and codecs,” in *2001 IEEE international conference on acoustics, speech, and signal processing. Proceedings (Cat. No. 01CH37221)*, vol. 2. IEEE, 2001, pp. 749–752.
- [7] J. G. Beerends, C. Schmidmer, J. Berger, M. Obermann, R. Ullmann, J. Pomy, and M. Keyhl, “Perceptual objective listening quality assessment (polqa), the third generation itu-t standard for end-to-end speech quality measurement part i—temporal alignment,” *Journal of the Audio Engineering Society*, vol. 61, no. 6, pp. 366–384, 2013.
- [8] A. Hines, J. Skoglund, A. Kokaram, and N. Harte, “Visqol: The virtual speech quality objective listener,” in *IWAENC 2012; International Workshop on Acoustic Signal Enhancement*. VDE, 2012, pp. 1–4.
- [9] M. Chinen, F. S. Lim, J. Skoglund, N. Gureev, F. O’Gorman, and A. Hines, “Visqol v3: An open source production ready objective speech and audio metric,” pp. 1–6, 2020.
- [10] J. Kim, M. El-Kharmy, and J. Lee, “End-to-end multi-task denoising for joint sdr and pesq optimization,” *arXiv preprint arXiv:1901.09146*, 2019.
- [11] S. Van Kuyk, W. B. Kleijn, and R. C. Hendriks, “An instrumental intelligibility metric based on information theory,” *IEEE Signal Processing Letters*, vol. 25, no. 1, pp. 115–119, 2017.

- [12] T. Thiede, W. C. Treurniet, R. Bitto, C. Schmidmer, T. Sporer, J. G. Beerends, and C. Colomes, “Peaq-the itu standard for objective measurement of perceived audio quality,” *Journal of the Audio Engineering Society*, vol. 48, no. 1/2, pp. 3–29, 2000.
- [13] A. Hines, E. Gillen, D. Kelly, J. Skoglund, A. Kokaram, and N. Harte, “Visqo-laudio: An objective audio quality metric for low bitrate codecs,” *The Journal of the Acoustical Society of America*, vol. 137, no. 6, pp. EL449–EL455, 2015.
- [14] J. Francombe, R. Mason, M. Dewhurst, and S. Bech, “A model of distraction in an audio-on-audio interference situation with music program material,” *Journal of the Audio Engineering Society*, vol. 63, no. 1/2, pp. 63–77, 2015.
- [15] —, “Elicitation of attributes for the evaluation of audio-on-audio interference,” *The Journal of the Acoustical Society of America*, vol. 136, no. 5, pp. 2630–2641, 2014.
- [16] J. Rämö, S. Bech, and S. H. Jensen, “Real-time perceptual model for distraction in interfering audio-on-audio scenarios,” *IEEE Signal Processing Letters*, vol. 24, no. 10, pp. 1448–1452, 2017.
- [17] C. H. Taal, R. C. Hendriks, and R. Heusdens, “A low-complexity spectro-temporal distortion measure for audio processing applications,” *IEEE Transactions on Audio, Speech, and Language Processing*, vol. 20, no. 5, pp. 1553–1564, 2012.
- [18] I. J. S. 29, “Information technology — coding of moving pictures and associated audio for digital storage media at up to about 1,5 mbit/s — part 3: Audio,” International Organization for Standardization, Geneva, CH, techreport 3, Oct. 1993.
- [19] D. Pan, “A tutorial on mpeg/audio compression,” *IEEE multimedia*, vol. 2, no. 2, pp. 60–74, 1995.
- [20] S. van de Par, A. Kohlrausch, R. Heusdens, J. Jensen, and S. H. Jensen, “A perceptual model for sinusoidal audio coding based on spectral integration,” *EURASIP Journal on Advances in Signal Processing*, vol. 2005, no. 9, pp. 1–13, 2005.
- [21] P. Balazs, B. Laback, G. Eckel, and W. A. Deutsch, “Time–frequency sparsity by removing perceptually irrelevant components using a simple model of simultaneous masking,” *IEEE transactions on audio, speech, and language processing*, vol. 18, no. 1, pp. 34–49, 2009.
- [22] C. H. Taal, J. Jensen, and A. Leijon, “On optimal linear filtering of speech for near-end listening enhancement,” *IEEE Signal Processing Letters*, vol. 20, no. 3, pp. 225–228, 2013.
- [23] S. Boyd, S. P. Boyd, and L. Vandenberghe, *Convex optimization*. Cambridge university press, 2004.
- [24] T. Painter and A. Spanias, “Perceptual coding of digital audio,” *Proceedings of the IEEE*, vol. 88, no. 4, pp. 451–515, 2000.

- [25] M. Olik, J. Francombe, P. Coleman, P. J. Jackson, M. Olsen, M. Møller, R. Mason, and S. Bech, “A comparative performance study of sound zoning methods in a reflective environment,” in *Audio Engineering Society Conference: 52nd International Conference: Sound Field Control-Engineering and Perception*. Audio Engineering Society, 2013.
- [26] S. J. Elliott and J. Cheer, “Regularisation and robustness of personal audio systems,” 2011.
- [27] Y. Cai, M. Wu, L. Liu, and J. Yang, “Time-domain acoustic contrast control design with response differential constraint in personal audio systems,” *The Journal of the Acoustical Society of America*, vol. 135, no. 6, pp. EL252–EL257, 2014.
- [28] M. B. Møller and M. Olsen, “Sound zones: On performance prediction of contrast control methods,” in *Audio Engineering Society Conference: 2016 AES International Conference on Sound Field Control*. Audio Engineering Society, 2016.

Extra Results

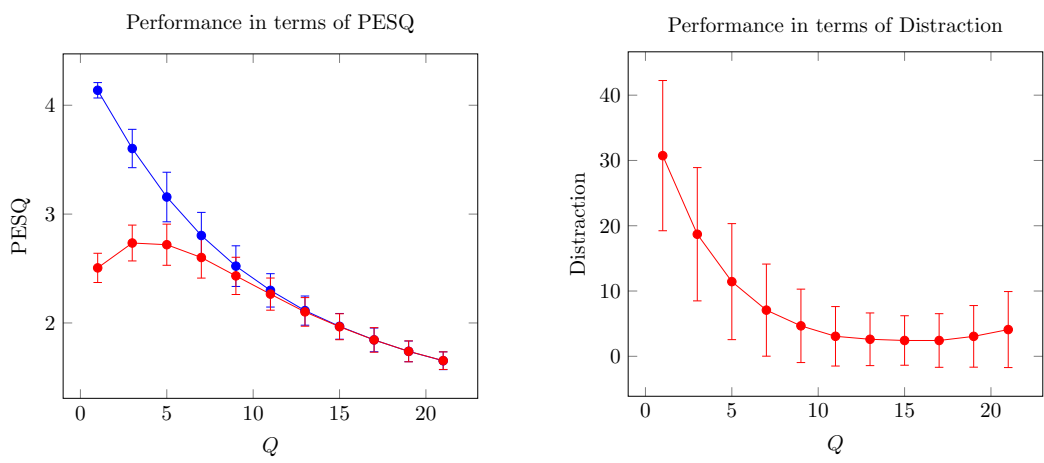


Figure A.1

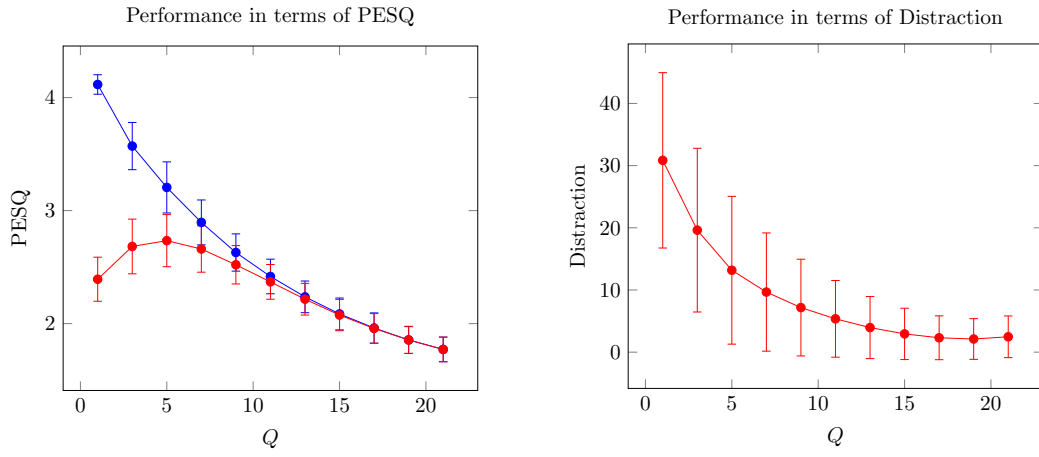


Figure A.2

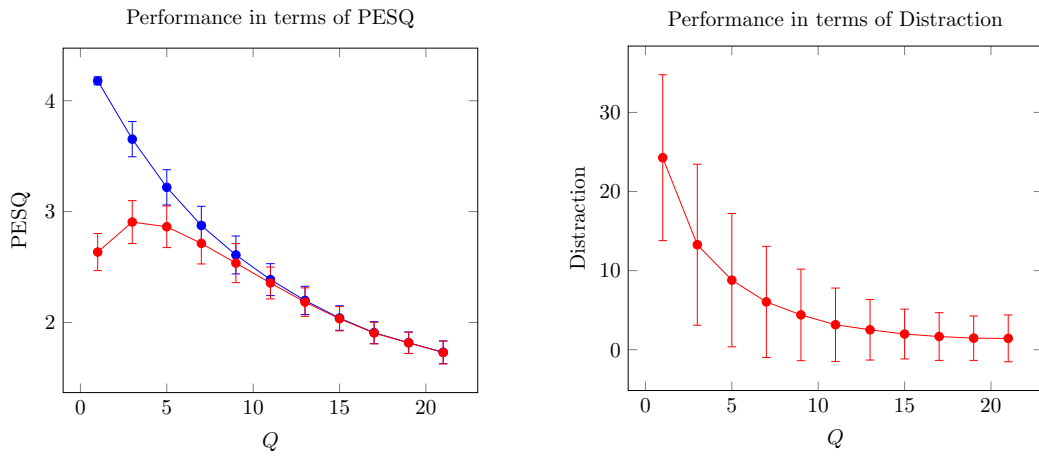


Figure A.3